

Importing Libraries

```
In [ ]: !pip install catboost
```

```
In [ ]: !pip uninstall -y scikit-learn
!pip install scikit-learn==1.3.1
```

```
In [ ]: !pip install imblearn
```

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix
from xgboost import XGBClassifier
from sklearn.svm import SVC
from catboost import CatBoostClassifier
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import cross_val_score
```

Reading The Dataset

```
In [5]: # Loading the dataset
file_path = './WA_Fn-UseC_-Telco-Customer-Churn.csv'

data = pd.read_csv(file_path)
print(f'dataset contains {data.shape[0]} rows and {data.shape[1]} columns')

dataset contains 7043 rows and 21 columns
```

```
In [6]: data.head(10)
```

```
Out[6]:
```

	customerID	gender	Senior_Citizen	Is_Married	Dependents	tenure	Phone_Service	Dual	Internet_Service	Online_Sec
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	
5	9305-CDSKC	Female	0	No	No	8	Yes	Yes	Fiber optic	
6	1452-KIOVK	Male	0	No	Yes	22	Yes	Yes	Fiber optic	
7	6713-OKOMC	Female	0	No	No	10	No	No phone service	DSL	
8	7892-POOKP	Female	0	Yes	No	28	Yes	Yes	Fiber optic	
9	6388-TABGU	Male	0	No	Yes	62	Yes	No	DSL	

10 rows × 21 columns

```
In [7]: data.tail()
```

Out[7]:	customerID	gender	Senior_Citizen	Is_Married	Dependents	tenure	Phone_Service	Dual	Internet_Service	Online_
7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	Yes	DSL	
7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	Yes	Fiber optic	
7040	4801-JZAZL	Female	0	Yes	Yes	11	No	No phone service	DSL	
7041	8361-LTMKD	Male	1	Yes	No	4	Yes	Yes	Fiber optic	
7042	3186-AJIEK	Male	0	No	No	66	Yes	No	Fiber optic	

5 rows × 21 columns

Get Information about the dataset:

```
In [8]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   Senior_Citizen        7043 non-null   int64
3   Is_Married            7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   Phone_Service         7043 non-null   object
7   Dual                  7043 non-null   object
8   Internet_Service      7043 non-null   object
9   Online_Security       7043 non-null   object
10  Online_Backup         7043 non-null   object
11  Device_Protection     7043 non-null   object
12  Tech_Support          7043 non-null   object
13  Streaming_TV          7043 non-null   object
14  Streaming_Movies      7043 non-null   object
15  Contract              7043 non-null   object
16  Paperless_Billing     7043 non-null   object
17  Payment_Method        7043 non-null   object
18  Monthly_Charges       7043 non-null   float64
19  Total_Charges         7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

Covertng Some of the Object Datatype Columns to Numerical

```
In [9]: # Total_charges column should be of numerical type
data['Total_Charges'] = data['Total_Charges'].apply(pd.to_numeric, errors='coerce') # invalid parsing will be s
```

Checking Null Values

```
In [10]: data.isnull().sum()
```

Out [10]:

	0
customerID	0
gender	0
Senior_Citizen	0
Is_Married	0
Dependents	0
tenure	0
Phone_Service	0
Dual	0
Internet_Service	0
Online_Security	0
Online_Backup	0
Device_Protection	0
Tech_Support	0
Streaming_TV	0
Streaming_Movies	0
Contract	0
Paperless_Billing	0
Payment_Method	0
Monthly_Charges	0
Total_Charges	11
Churn	0

dtype: int64

In [11]:

```
# Check rows where Total Charges is NaN
nan_rows = data[data['Total_Charges'].isna()]
print(nan_rows)
```

	customerID	gender	Senior_Citizen	Is_Married	Dependents	tenure	\
488	4472-LVYGI	Female	0	Yes	Yes	0	
753	3115-CZMZD	Male	0	No	Yes	0	
936	5709-LV0EQ	Female	0	Yes	Yes	0	
1082	4367-NUYAO	Male	0	Yes	Yes	0	
1340	1371-DWPAZ	Female	0	Yes	Yes	0	
3331	7644-OMVMY	Male	0	Yes	Yes	0	
3826	3213-VVOLG	Male	0	Yes	Yes	0	
4380	2520-SGTTA	Female	0	Yes	Yes	0	
5218	2923-ARZLG	Male	0	Yes	Yes	0	
6670	4075-WKNIU	Female	0	Yes	Yes	0	
6754	2775-SEFEE	Male	0	No	Yes	0	

	Phone_Service	Dual	Internet_Service	Online_Security	\
488	No	No phone service	DSL	Yes	
753	Yes	No	No	No internet service	
936	Yes	No	DSL	Yes	
1082	Yes	Yes	No	No internet service	
1340	No	No phone service	DSL	Yes	
3331	Yes	No	No	No internet service	
3826	Yes	Yes	No	No internet service	
4380	Yes	No	No	No internet service	
5218	Yes	No	No	No internet service	
6670	Yes	Yes	DSL	No	
6754	Yes	Yes	DSL	Yes	

	Device_Protection	Tech_Support	Streaming_TV	\
488	...	Yes	Yes	
753	...	No internet service	No internet service	
936	...	Yes	No	
1082	...	No internet service	No internet service	
1340	...	Yes	Yes	
3331	...	No internet service	No internet service	
3826	...	No internet service	No internet service	
4380	...	No internet service	No internet service	
5218	...	No internet service	No internet service	
6670	...	Yes	Yes	
6754	...	No	Yes	

	Streaming_Movies	Contract	Paperless_Billing	\
488	No	Two year	Yes	
753	No internet service	Two year	No	
936	Yes	Two year	No	
1082	No internet service	Two year	No	
1340	No	Two year	No	
3331	No internet service	Two year	No	
3826	No internet service	Two year	No	
4380	No internet service	Two year	No	
5218	No internet service	One year	Yes	
6670	No	Two year	No	
6754	No	Two year	Yes	

	Payment_Method	Monthly_Charges	Total_Charges	Churn
488	Bank transfer (automatic)	52.55	NaN	No
753	Mailed check	20.25	NaN	No
936	Mailed check	80.85	NaN	No
1082	Mailed check	25.75	NaN	No
1340	Credit card (automatic)	56.05	NaN	No
3331	Mailed check	19.85	NaN	No
3826	Mailed check	25.35	NaN	No
4380	Mailed check	20.00	NaN	No
5218	Mailed check	19.70	NaN	No
6670	Mailed check	73.35	NaN	No
6754	Bank transfer (automatic)	61.90	NaN	No

[11 rows x 21 columns]

```
In [12]: # Due to the dataset being big i decided to drop the rows
```

```
# Drop rows with NaN in the 'Total_Charges' column
data = data.dropna(subset=['Total_Charges'])
```

```
In [13]: data['Total_Charges'].isnull().sum()
```

```
Out[13]: 0
```

Checking Duplicates

```
In [14]: data_dup = data.duplicated().any()
print(data_dup)
# data.drop_duplicates()
```

```
False
```

Overall statistics about the dataset

```
In [15]: data.describe()
```

Out[15]:

	Senior_Citizen	tenure	Monthly_Charges	Total_Charges
count	7032.000000	7032.000000	7032.000000	7032.000000
mean	0.162400	32.421786	64.798208	2283.300441
std	0.368844	24.545260	30.085974	2266.771362
min	0.000000	1.000000	18.250000	18.800000
25%	0.000000	9.000000	35.587500	401.450000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.862500	3794.737500
max	1.000000	72.000000	118.750000	8684.800000

Splitting the Dataset into Features and a Target

```
In [16]: X = data.iloc[:, 1:-1] # Excluding the customerid and the churn
y = data.iloc[:, -1] # Churn
```

```
In [17]: print(X.shape)
print(y.shape)

(7032, 19)
(7032,)
```

```
In [18]: X.head(5)
```

Out[18]:

	gender	Senior_Citizen	Is_Married	Dependents	tenure	Phone_Service	Dual	Internet_Service	Online_Security	Online
0	Female	0	Yes	No	1	No	No phone service	DSL	No	
1	Male	0	No	No	34	Yes	No	DSL	Yes	
2	Male	0	No	No	2	Yes	No	DSL	Yes	
3	Male	0	No	No	45	No	No phone service	DSL	Yes	
4	Female	0	No	No	2	Yes	No	Fiber optic	No	

```
In [19]: y.head(5)
```

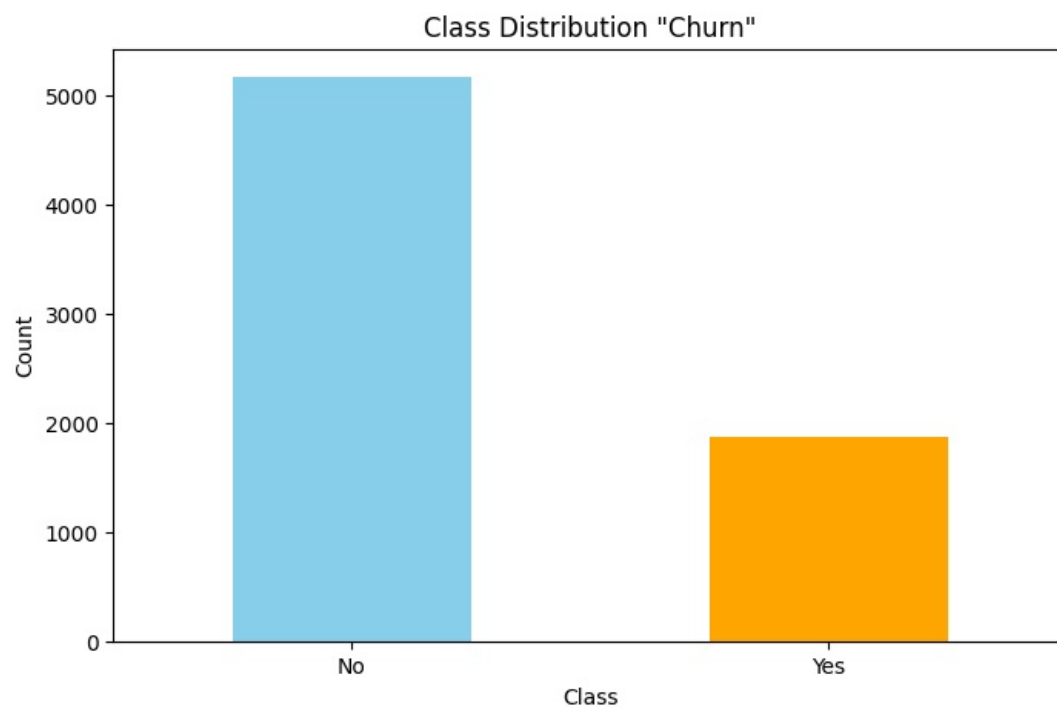
Out[19]:

	Churn
0	No
1	No
2	Yes
3	No
4	Yes


dtype: object

```
In [20]: # Get the count of each class
class_counts = y.value_counts() # Counts of unique values

# Plot the counts as a bar chart
plt.figure(figsize=(8, 5))
class_counts.plot(kind='bar', color=['skyblue', 'orange'])
plt.title('Class Distribution "Churn"')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.show()
```



```
In [21]: class_counts.plot.pie(autopct='%1.2f%%', colors=['skyblue', 'orange'])
plt.title('Class Distribution "Churn"')
plt.ylabel('')
plt.show()
# Looks Like an Imbalanced Dataset
```



Encoding Categorical Data

```
In [22]: # Checking unique values to choose which technique to apply
should_be_one_hot_encoded = []
should_be_label_encoded = []

for col in X.columns:
    if X[col].dtypes == 'object': # Exclude numerical values
        print(f'{col}: {X[col].unique()}')
        if len(X[col].unique()) > 2:
            should_be_one_hot_encoded.append(col)
        else:
            should_be_label_encoded.append(col)

print('\nOne-Hot Encoded : ', should_be_one_hot_encoded, '\n')
print('Label Encoded : ', should_be_label_encoded)
```

```

gender: ['Female' 'Male']
Is_Married: ['Yes' 'No']
Dependents: ['No' 'Yes']
Phone_Service: ['No' 'Yes']
Dual: ['No phone service' 'No' 'Yes']
Internet_Service: ['DSL' 'Fiber optic' 'No']
Online_Security: ['No' 'Yes' 'No internet service']
Online_Backup: ['Yes' 'No' 'No internet service']
Device_Protection: ['No' 'Yes' 'No internet service']
Tech_Support: ['No' 'Yes' 'No internet service']
Streaming_TV: ['No' 'Yes' 'No internet service']
Streaming_Movies: ['No' 'Yes' 'No internet service']
Contract: ['Month-to-month' 'One year' 'Two year']
Paperless_Billing: ['Yes' 'No']
Payment_Method: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']

```

```

One-Hot Encoded : ['Dual', 'Internet_Service', 'Online_Security', 'Online_Backup', 'Device_Protection', 'Tech_Support', 'Streaming_TV', 'Streaming_Movies', 'Contract', 'Payment_Method']

```

```

Label Encoded : ['gender', 'Is_Married', 'Dependents', 'Phone_Service', 'Paperless_Billing']

```

Apply Label Encoding

```

In [23]: le = LabelEncoder()
for col in should_be_label_encoded:
    X[col] = le.fit_transform(X[col]) # Apply label encoding for each column

for col in should_be_label_encoded:
    print(f'{col}: {X[col].unique()}')

```

```

gender: [0 1]
Is_Married: [1 0]
Dependents: [0 1]
Phone_Service: [0 1]
Paperless_Billing: [1 0]

```

```

In [24]: # Label Encoding the Target
y = le.fit_transform(y)
print(y)

```

```

[0 0 1 ... 0 1 0]

```

Apply One-Hot Encoding

```

In [25]: # Get the indexes of columns to transform
hot_encode_indexes = X.columns.get_indexer(should_be_one_hot_encoded)
print(hot_encode_indexes)
ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), hot_encode_indexes)], remainder='passthrough')
# Fit and transform the data
X_transformed = np.array(ct.fit_transform(X))
print(X_transformed)

```

```

[ 6  7  8  9 10 11 12 13 14 16]
[[0.0000e+00 1.0000e+00 0.0000e+00 ... 1.0000e+00 2.9850e+01 2.9850e+01]
 [1.0000e+00 0.0000e+00 0.0000e+00 ... 0.0000e+00 5.6950e+01 1.8895e+03]
 [1.0000e+00 0.0000e+00 0.0000e+00 ... 1.0000e+00 5.3850e+01 1.0815e+02]
 ...
 [0.0000e+00 1.0000e+00 0.0000e+00 ... 1.0000e+00 2.9600e+01 3.4645e+02]
 [0.0000e+00 0.0000e+00 1.0000e+00 ... 1.0000e+00 7.4400e+01 3.0660e+02]
 [1.0000e+00 0.0000e+00 0.0000e+00 ... 1.0000e+00 1.0565e+02 6.8445e+03]]

```

```

In [26]: # Get the feature names for the one-hot encoded columns
encoder = ct.transformers_[0][1] # The encoder used for one-hot encoding
encoded_feature_names = encoder.get_feature_names_out(input_features=should_be_one_hot_encoded)

# Create a DataFrame with the transformed data
# Concatenate the new one-hot encoded feature names and original columns that weren't transformed
X_transformed_df = pd.DataFrame(X_transformed, columns=np.concatenate([encoded_feature_names, X.columns.difference(encoded_feature_names)]))

# Show the resulting DataFrame
print(X_transformed_df)

```

	Dual_No	Dual_No phone service	Dual_Yes	Internet_Service_DSL	\
0	0.0	1.0	0.0	1.0	
1	1.0	0.0	0.0	1.0	
2	1.0	0.0	0.0	1.0	
3	0.0	1.0	0.0	1.0	
4	1.0	0.0	0.0	0.0	
...	
7027	0.0	0.0	1.0	1.0	
7028	0.0	0.0	1.0	0.0	
7029	0.0	1.0	0.0	1.0	
7030	0.0	0.0	1.0	0.0	
7031	1.0	0.0	0.0	0.0	

	Internet_Service_Fiber optic	Internet_Service_No	Online_Security_No	\
0	0.0	0.0	1.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	1.0	0.0	1.0	
...	
7027	0.0	0.0	0.0	
7028	1.0	0.0	1.0	
7029	0.0	0.0	0.0	
7030	1.0	0.0	1.0	
7031	1.0	0.0	0.0	

	Online_Security_No internet service	Online_Security_Yes	\
0	0.0	0.0	
1	0.0	1.0	
2	0.0	1.0	
3	0.0	1.0	
4	0.0	0.0	
...	
7027	0.0	1.0	
7028	0.0	0.0	
7029	0.0	1.0	
7030	0.0	0.0	
7031	0.0	1.0	

	Online_Backup_No	...	Payment_Method_Mailed check	Dependents	\
0	0.0	...	0.0	0.0	
1	1.0	...	1.0	1.0	
2	0.0	...	1.0	1.0	
3	1.0	...	0.0	1.0	
4	1.0	...	0.0	0.0	
...	
7027	1.0	...	1.0	1.0	
7028	0.0	...	0.0	0.0	
7029	1.0	...	0.0	0.0	
7030	1.0	...	1.0	1.0	
7031	1.0	...	0.0	1.0	

	Is_Married	Monthly_Charges	Paperless_Billing	Phone_Service	\
0	0.0	1.0	0.0	1.0	
1	0.0	0.0	0.0	34.0	
2	0.0	0.0	0.0	2.0	
3	0.0	0.0	0.0	45.0	
4	0.0	0.0	0.0	2.0	
...	
7027	0.0	1.0	1.0	24.0	
7028	0.0	1.0	1.0	72.0	
7029	0.0	1.0	1.0	11.0	
7030	1.0	1.0	0.0	4.0	
7031	0.0	0.0	0.0	66.0	

	Senior_Citizen	Total_Charges	gender	tenure
0	0.0	1.0	29.85	29.85
1	1.0	0.0	56.95	1889.50
2	1.0	1.0	53.85	108.15
3	0.0	0.0	42.30	1840.75
4	1.0	1.0	70.70	151.65
...
7027	1.0	1.0	84.80	1990.50
7028	1.0	1.0	103.20	7362.90
7029	0.0	1.0	29.60	346.45
7030	1.0	1.0	74.40	306.60
7031	1.0	1.0	105.65	6844.50

[7032 rows x 40 columns]

Feature Selection Using Correlation

```
In [27]: #Get Correlation of "Churn" with other variables:
plt.figure(figsize=(15,8))
y_df = pd.DataFrame(y, columns=['Churn'])

# Concatenate the feature DataFrame and the target DataFrame
```




```
new_df = pd.concat([X_transformed_df, y_df], axis=1)

# Calculate correlations with the 'Churn' column
correlation = new_df.corr()['Churn'].sort_values(ascending=False)

# Plot the correlation of Churn with other variables
plt.figure(figsize=(15, 8))
correlation.plot(kind='bar')
plt.title("Correlation of Features with Churn")
plt.ylabel("Correlation Coefficient")
plt.show()
```

<Figure size 1500x800 with 0 Axes>



Feature Selection Using RandomForest

```
In [28]: # Dataset has mixed values between numerical and categorical so its best to use randomforest rather than correl
model = RandomForestClassifier()
model.fit(X_transformed, y)


# Get feature importances
importances = model.feature_importances_
print(importances)

# Create a DataFrame to store feature names and their importances
features_df = pd.DataFrame({
    'Feature': X_transformed_df.columns,
    'Importance': importances
})

# Sort the features by importance in descending order
features_df = features_df.sort_values(by='Importance', ascending=True)

# Plot the feature importances
plt.figure(figsize=(15, 8))
plt.barh(features_df['Feature'], features_df['Importance'], color='purple')
plt.xlabel('Importance')
plt.title('Feature Importances from RandomForestClassifier')
plt.show()
```

```
[0.0134645 0.00362119 0.01211903 0.01114978 0.02472343 0.00236605
0.02996051 0.00290799 0.01045103 0.0183466 0.00095171 0.0115077
0.01623842 0.00317968 0.01109489 0.03035744 0.00331088 0.01148571
0.01093526 0.00247654 0.01134598 0.01143882 0.00173175 0.01090508
0.05363645 0.00967561 0.01567628 0.01177007 0.01208563 0.02799879
0.01081188 0.02810681 0.02097648 0.02340859 0.01894302 0.15290328
0.00322414 0.02444555 0.14982801 0.17043944]
```



Keeping High-Correlated Features Only

```
In [29]: number_of_features = 15 # Take highest 15 features
filtered_features = []
for feature in features_df.tail(number_of_features).Feature:
    print(feature)
    filtered_features.append(feature)

for feature in X_transformed_df.columns:
    if feature not in filtered_features:
        X_transformed_df = X_transformed_df.drop(feature, axis=1)

print(X_transformed_df)
```

```
Online_Security_Yes
Payment_Method_Mailed check
Streaming_Movies_Yes
Streaming_TV_No
Device_Protection_Yes
Internet_Service_DSL
Streaming_TV_Yes
Streaming_Movies_No
Tech_Support_Yes
Online_Backup_Yes
Payment_Method_Bank transfer (automatic)
Payment_Method_Credit card (automatic)
Dual_Yes
Dual_No
Contract_Two year
Device_Protection_No
Online_Backup_No
Paperless_Billing
Is_Married
Monthly_Charges
Total_Charges
Internet_Service_Fiber optic
Payment_Method_Electronic check
Dependents
Online_Security_No
Tech_Support_No
Contract_Month-to-month
gender
Phone_Service
tenure
```

	Dual_No	Dual_Yes	Internet_Service_DSL	Internet_Service_Fiber optic	\
0	0.0	0.0	1.0	0.0	
1	1.0	0.0	1.0	0.0	
2	1.0	0.0	1.0	0.0	
3	0.0	0.0	1.0	0.0	
4	1.0	0.0	0.0	1.0	
...	
7027	0.0	1.0	1.0	0.0	
7028	0.0	1.0	0.0	1.0	

7029	0.0	0.0	1.0	0.0
7030	0.0	1.0	0.0	1.0
7031	1.0	0.0	0.0	1.0

	Online_Security_No	Online_Security_Yes	Online_Backup_No	\
0	1.0	0.0	0.0	
1	0.0	1.0	1.0	
2	0.0	1.0	0.0	
3	0.0	1.0	1.0	
4	1.0	0.0	1.0	
...	
7027	0.0	1.0	1.0	
7028	1.0	0.0	0.0	
7029	0.0	1.0	1.0	
7030	1.0	0.0	1.0	
7031	0.0	1.0	1.0	

	Online_Backup_Yes	Device_Protection_No	Device_Protection_Yes	...	\
0	1.0	1.0	0.0	...	
1	0.0	0.0	1.0	...	
2	1.0	1.0	0.0	...	
3	0.0	0.0	1.0	...	
4	0.0	1.0	0.0	...	
...	
7027	0.0	0.0	1.0	...	
7028	1.0	0.0	1.0	...	
7029	0.0	1.0	0.0	...	
7030	0.0	1.0	0.0	...	
7031	0.0	0.0	1.0	...	

	Payment_Method_Electronic	check	Payment_Method_Mailed	check	\
0		1.0		0.0	
1		0.0		1.0	
2		0.0		1.0	
3		0.0		0.0	
4		1.0		0.0	
...		
7027		0.0		1.0	
7028		0.0		0.0	
7029		1.0		0.0	
7030		0.0		1.0	
7031		0.0		0.0	

	Dependents	Is_Married	Monthly_Charges	Paperless_Billing	\
0	0.0	0.0	1.0	0.0	
1	1.0	0.0	0.0	0.0	
2	1.0	0.0	0.0	0.0	
3	1.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	
...	
7027	1.0	0.0	1.0	1.0	
7028	0.0	0.0	1.0	1.0	
7029	0.0	0.0	1.0	1.0	
7030	1.0	1.0	1.0	0.0	
7031	1.0	0.0	0.0	0.0	

	Phone_Service	Total_Charges	gender	tenure
0	1.0	1.0	29.85	29.85
1	34.0	0.0	56.95	1889.50
2	2.0	1.0	53.85	108.15
3	45.0	0.0	42.30	1840.75
4	2.0	1.0	70.70	151.65
...
7027	24.0	1.0	84.80	1990.50
7028	72.0	1.0	103.20	7362.90
7029	11.0	1.0	29.60	346.45
7030	4.0	1.0	74.40	306.60
7031	66.0	1.0	105.65	6844.50

[7032 rows x 30 columns]

Splitting the dataset into the Training and Test sets

```
In [30]: X_train, X_test, y_train, y_test = train_test_split(np.array(X_transformed_df), y, test_size=0.2, random_state=
```

Feature Scaling

```
In [31]: sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

Perform UnderSampling

```
# from sklearn under sampling import RandomUnderSampler
```

```
In [32]: # from imblearn.under_sampling import RandomUnderSampler
# rus = RandomUnderSampler(random_state=42)
# X_train, y_train = rus.fit_resample(X_train, y_train)
```


Perform OverSampling

```
In [33]: # from imblearn.over_sampling import SMOTE
# smote = SMOTE(random_state=42)
# X_train, y_train = smote.fit_resample(X_train, y_train)
```

Perform UnderOverSampling

```
In [34]: from imblearn.combine import SMOTEENN
smote = SMOTEENN(sampling_strategy='auto', random_state=42)
X_train, y_train = smote.fit_resample(X_train, y_train)
```

```
In [35]: pd.DataFrame(y_train).value_counts().plot.pie(autopct='%1.2f%%', colors=['skyblue', 'orange'])
plt.title('Class Distribution "Churn"')
plt.ylabel('')
plt.show()
```



Training XGBoost on the Training set

```
In [36]: xgb_classifier = XGBClassifier(n_estimators=194, learning_rate=0.08011727992081952, max_depth=3)
# {'n_estimators': 138, 'learning_rate': 0.2910194858724226, 'max_depth': 8}

xgb_classifier.fit(X_train, y_train)
```

```
Out[36]: XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.0801172799208195
2,
              max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=3, max_leaves=None,
```

Confusion Matrix

```
In [37]: y_pred = xgb_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[719 314]
 [ 80 294]]
0.7199715707178393
```

```
Out[37]:
```

K-Cross Validation

```
In [38]: accuracies = cross_val_score(estimator = xgb_classifier, X = X_train, y = y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the variance
```

Accuracy: 95.65 %
Standard Deviation: 2.34 %

```
In [39]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.70	0.78	1033
1	0.48	0.79	0.60	374
accuracy			0.72	1407
macro avg	0.69	0.74	0.69	1407
weighted avg	0.79	0.72	0.74	1407

```
In [ ]: !pip install optuna
```

Hyper Parameter Tuning

```
In [ ]: import optuna
# from sklearn.metrics import accuracy_score

def objective(trial):
    params = {
        'n_estimators': trial.suggest_int('n_estimators', 50, 200),
        'learning_rate': trial.suggest_loguniform('learning_rate', 0.01, 0.3),
        'max_depth': trial.suggest_int('max_depth', 3, 9),
    }
    model = XGBClassifier(**params)
    # model.fit(X_train, y_train)
    # Evaluate on the test set
    # test_accuracy = accuracy_score(y_test, model.predict(X_test))
    # return test_accuracy
    return cross_val_score(model, X_train, y_train, cv=5, scoring='accuracy').mean()
study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=50)
print(study.best_params)
```

Trying Randomized Search

```
In [42]: # param_dist = {
#         'learning_rate': [0.01, 0.1, 0.2, 0.3],
#         'n_estimators': [50, 100, 200],
#         'max_depth': [3, 5, 7, 9],
#         'min_child_weight': [1, 3, 5],
#         'subsample': [0.6, 0.8, 1.0]
#     }

# random_search = RandomizedSearchCV(xgb_classifier, param_distributions=param_dist, n_iter=10, cv=3)
# random_search.fit(X_train, y_train)
# best_accuracy = random_search.best_score_
# best_parameters = random_search.best_params_
# print(f"Best Accuracy: {round(best_accuracy*100, 2)} %")
# print(f"Best Parameters: {best_parameters}")
```

Training on CatBoost

```
In [43]: cat_classifier = CatBoostClassifier(silent=True)
cat_classifier.fit(X_train, y_train)
```

```
Out[43]: <catboost.core.CatBoostClassifier at 0x7db287555e10>
```

```
In [44]: y_pred = cat_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[732 301]
 [ 83 291]]
0.7270788912579957
```

```
Out[44]:
```

```
In [45]: accuracies = cross_val_score(estimator = cat_classifier, X = X_train, y = y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
```

```
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the variance
```

Accuracy: 96.63 %
Standard Deviation: 1.93 %

```
In [46]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.71	0.79	1033
1	0.49	0.78	0.60	374
accuracy			0.73	1407
macro avg	0.69	0.74	0.70	1407
weighted avg	0.79	0.73	0.74	1407

```
In [47]: # import optuna
# def objective(trial):
#     # Suggest values for the hyperparameters
#     iterations = trial.suggest_categorical('iterations', [500, 1000])
#     learning_rate = trial.suggest_loguniform('learning_rate', 0.01, 0.2)
#     depth = trial.suggest_int('depth', 6, 10)
#     l2_leaf_reg = trial.suggest_float('l2_leaf_reg', 1, 5)

#     cat_classifier = CatBoostClassifier(
#         iterations=iterations,
#         learning_rate=learning_rate,
#         depth=depth,
#         l2_leaf_reg=l2_leaf_reg,
#         verbose=0
#     )
#     return cross_val_score(cat_classifier, X_train, y_train, cv=5, scoring='accuracy').mean()
# study = optuna.create_study(direction='maximize')
# study.optimize(objective, n_trials=50)
# print(study.best_params)
```

```
In [48]: # param_dist = {
#         #     'iterations': [500, 1000],
#         #     'learning_rate': [0.01, 0.1, 0.2],
#         #     'depth': [6, 8, 10],
#         #     'l2_leaf_reg': [1, 3, 5]
#         # }

# random_search = RandomizedSearchCV(cat_classifier, param_distributions=param_dist, n_iter=10, cv=3)
# random_search.fit(X_train, y_train)
# best_accuracy = random_search.best_score_
# best_parameters = random_search.best_params_
# print(f"Best Accuracy: {round(best_accuracy*100, 2)} %")
# print(f"Best Parameters: {best_parameters}")
```

Training KernelSVM on the Training set

```
In [49]: svm_classifier = SVC(kernel = 'rbf', random_state = 0)
svm_classifier.fit(X_train, y_train)
```

```
Out[49]: SVC
SVC(random_state=0)
```

```
In [50]: y_pred = svm_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[693 340]
 [ 80 294]]
Out[50]: 0.7014925373134329
```

```
In [51]: from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = svm_classifier, X = X_train, y = y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the variance
```

Accuracy: 94.29 %
Standard Deviation: 0.9 %

```
In [52]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.67	0.77	1033
1	0.46	0.79	0.58	374
accuracy			0.70	1407
macro avg	0.68	0.73	0.68	1407
weighted avg	0.78	0.70	0.72	1407

```
In [53]: # from sklearn.model_selection import GridSearchCV
# parameters = [{'C': [0.25, 0.5, 0.75, 1], 'kernel': ['linear']},
#               {'C': [0.25, 0.5, 0.75, 1], 'kernel': ['rbf'], 'gamma': [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]}]
# grid_search = GridSearchCV(estimator = svm_classifier,
#                             param_grid = parameters,
#                             scoring = 'accuracy', # the metric with which you want to evaluate the performance
#                             cv = 10, # number of trained test polls when applying k-fold cross validation
#                             n_jobs = -1) # number of cores to use
# grid_search.fit(X_train, y_train)
# best_accuracy = grid_search.best_score_
# best_parameters = grid_search.best_params_
# print(f"Best Accuracy: {round(best_accuracy*100, 2)} %")
# print(f"Best Parameters: {best_parameters}")
```

```
In [54]: # def objective(trial):
#         # Suggest values for the hyperparameters
#         c = trial.suggest_loguniform('C', 1e-3, 1e3) # Loguniform for C to cover a wide range
#         kernel = trial.suggest_categorical('kernel', ['linear', 'poly', 'rbf', 'sigmoid'])
#         degree = trial.suggest_int('degree', 2, 5) if kernel == 'poly' else 3 # Degree only relevant for 'poly'
#         gamma = trial.suggest_categorical('gamma', ['scale', 'auto'])

#         # Create an SVC with suggested parameters
#         svc = SVC(C=c, kernel=kernel, degree=degree, gamma=gamma)

#         # Perform cross-validation and return the mean accuracy
#         score = cross_val_score(svc, X_train, y_train, cv=5, scoring='accuracy').mean()
#         return score

# # Create and optimize the study
# study = optuna.create_study(direction='maximize')
# study.optimize(objective, n_trials=50)
# print(study.best_params_)
```

Training Logistic Regression on the Training set

```
In [55]: from sklearn.linear_model import LogisticRegression
lr_classifier = LogisticRegression(random_state = 0)
lr_classifier.fit(X_train, y_train)
```

```
Out[55]: LogisticRegression
LogisticRegression(random_state=0)
```

```
In [56]: y_pred = lr_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[651 382]
 [ 54 320]]
0.6901208244491827
```

```
In [57]: from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = lr_classifier, X = X_train, y = y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the variance
```

Accuracy: 90.09 %
Standard Deviation: 1.16 %

```
In [58]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.92	0.63	0.75	1033
1	0.46	0.86	0.59	374
accuracy			0.69	1407
macro avg	0.69	0.74	0.67	1407
weighted avg	0.80	0.69	0.71	1407

```
In [59]: # def objective(trial):
#         # Suggest values for the hyperparameters
```

```
# c = trial.suggest_loguniform('C', 1e-3, 1e3) # Regularization strength (log-uniform)
# solver = trial.suggest_categorical('solver', ['liblinear', 'saga', 'lbfgs', 'newton-cg'])
# max_iter = trial.suggest_int('max_iter', 100, 500) # Maximum iterations
# penalty = trial.suggest_categorical('penalty', ['l2', 'elasticnet']) # Penalty type

# # ElasticNet mixing parameter (only relevant if 'penalty' is 'elasticnet')
# l1_ratio = trial.suggest_float('l1_ratio', 0.0, 1.0) if penalty == 'elasticnet' else 0.5

# # Create Logistic Regression with suggested parameters
# logreg = LogisticRegression(
#     C=c,
#     solver=solver,
#     max_iter=max_iter,
#     penalty=penalty,
#     l1_ratio=l1_ratio,
#     multi_class='auto',
#     random_state=42
# )

# # Perform cross-validation and return the mean accuracy
# score = cross_val_score(logreg, X_train, y_train, cv=5, scoring='accuracy').mean()
# return score

# # Create and optimize the study
# study = optuna.create_study(direction='maximize')
# study.optimize(objective, n_trials=50)
# print(study.best_params)
```

Training Random Forest on the Training set

```
In [60]: from sklearn.ensemble import RandomForestClassifier
rf_classifier = RandomForestClassifier(n_estimators=220, max_depth=17, criterion='entropy', random_state=0,
rf_classifier.fit(X_train, y_train)
# {'n_estimators': 220, 'max_depth': 17, 'min_samples_split': 5, 'min_samples_leaf': 1, 'max_features': 'sqrt'}
```

```
Out[60]: ▼ Random Forest Classifier
RandomForestClassifier(criterion='entropy', max_depth=17, min_samples_split=5,
n_estimators=220, random_state=0)
```

```
In [61]: y_pred = rf_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[717 316]
 [ 81 293]]
Out[61]: 0.7178393745557925
```

```
In [62]: from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = rf_classifier, X = X_train, y = y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the variance
```

```
Accuracy: 96.43 %
Standard Deviation: 1.21 %
```

```
In [63]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.69	0.78	1033
1	0.48	0.78	0.60	374
accuracy			0.72	1407
macro avg	0.69	0.74	0.69	1407
weighted avg	0.79	0.72	0.73	1407

```
In [64]: # import optuna
# #{'n_estimators': 220, 'max_depth': 17, 'min_samples_split': 5, 'min_samples_leaf': 1, 'max_features': 'sqrt'}
# def objective(trial):
#     # Suggest values for the hyperparameters
#     n_estimators = trial.suggest_int('n_estimators', 50, 500) # Number of trees
#     max_depth = trial.suggest_int('max_depth', 2, 20) # Maximum depth of trees
#     min_samples_split = trial.suggest_int('min_samples_split', 2, 20) # Minimum samples to split a node
#     min_samples_leaf = trial.suggest_int('min_samples_leaf', 1, 20) # Minimum samples per leaf
#     max_features = trial.suggest_categorical('max_features', ['sqrt', 'log2', None]) # Number of features to

# # Create a RandomForestClassifier with suggested parameters
# rf = RandomForestClassifier(
#     n_estimators=n_estimators,
#     max_depth=max_depth,
#     min_samples_split=min_samples_split,
```



```
#         min_samples_leaf=min_samples_leaf,
#         max_features=max_features,
#         random_state=42
#     )

#     # Perform cross-validation and return the mean accuracy
#     score = cross_val_score(rf, X_train, y_train, cv=5, scoring='accuracy').mean()
#     return score

# # Create and optimize the study
# study = optuna.create_study(direction='maximize')
# study.optimize(objective, n_trials=50)
# print(study.best_params)
```

Building an ANN

```
In [76]: import tensorflow as tf
from tensorflow.keras.regularizers import l2
# Building the Network

model = tf.keras.models.Sequential(
    [
        tf.keras.Input(shape=(X_train.shape[1],)),
        tf.keras.layers.Dense(units=256, activation='relu'),
        tf.keras.layers.BatchNormalization(),
        tf.keras.layers.Dropout(0.3),
        tf.keras.layers.Dense(units=256, activation='relu', kernel_regularizer=l2(0.001)),
        tf.keras.layers.Dropout(0.2),
        tf.keras.layers.Dense(units=128, activation='relu'),
        tf.keras.layers.Dense(units=1, activation='sigmoid')
    ]
)
```

```
In [77]: def plot_loss_acc(history):
    '''Plots the training and validation loss and accuracy from a history object'''
    acc = history.history['accuracy']
    loss = history.history['loss']
    val_acc = history.history['val_accuracy']
    val_loss = history.history['val_loss']

    epochs = range(len(acc))

    fig, ax = plt.subplots(1,2, figsize=(12, 6))
    ax[0].plot(epochs, acc, 'b', label='Training accuracy')
    ax[0].plot(epochs, val_acc, 'r', label='Validation accuracy')
    ax[0].set_title('Training accuracy')
    ax[0].set_xlabel('epochs')
    ax[0].set_ylabel('accuracy')
    ax[0].legend()

    ax[1].plot(epochs, loss, 'b', label='Training Loss')
    ax[1].plot(epochs, val_loss, 'r', label='Validation Loss')
    ax[1].set_title('Training loss')
    ax[1].set_xlabel('epochs')
    ax[1].set_ylabel('loss')
    ax[1].legend()

    plt.show()
```

```
In [78]: class EarlyStoppingCallback(tf.keras.callbacks.Callback):

    # Define the correct function signature for on_epoch_end method
    def on_epoch_end(self, epoch, logs={}):

        # Check if the accuracy is greater or equal to 0.98
        if (logs.get("accuracy") >= 0.94):

            # Stop training once the above condition is met
            self.model.stop_training = True

            print("\nReached 94% accuracy so cancelling training!")
```

```
In [79]: def plot_learning_rate(history):
    """
    Plot learning rate vs loss.
    """
    # Define the learning rate array
    lrs = 1e-8 * (10 ** (np.arange(len(history.history["loss"])) / 20))

    # Set the figure size
    plt.figure(figsize=(10, 6))

    # Set the grid
    plt.grid(True)

    # Plot the loss in log scale
```

```
plt.semilogx(lrs, history.history["loss"])

# Increase the tickmarks size
plt.tick_params('both', length=10, width=1, which='both')

# Set the plot boundaries
plt.axis([1e-8, 8e-4, 0, 50])
plt.ylim(0, 1.0) # This would zoom the y-axis between 0 and 1.0

# Draw the graph on screen
plt.xlabel('Learning Rate')
plt.ylabel('Loss')
plt.title('Learning Rate vs Loss')
plt.show()
```

```
In [80]: def tuning_learn_rate(model, X_train, y_train, optimizer, loss, batch_size=64, epochs=100):
        """
        Tunes the learning rate using an exponential increase.

        Args:
            model (tf.keras.Model): The Keras model.
            X_train (np.array): Training features.
            y_train (np.array): Training labels.
            optimizer (tf.keras.optimizers): Optimizer to use.
            loss (str or tf.keras.losses.Loss): Loss function for the model.
            epochs (int): Number of epochs to tune learning rate.
            batch_size (int): The batch size for training.

        Returns:
            history (tf.keras.callbacks.History): Training history object.
        """
        # Convert the data into a TensorFlow Dataset
        train_dataset = tf.data.Dataset.from_tensor_slices((X_train, y_train))
        train_dataset = train_dataset.batch(batch_size)




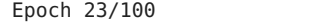
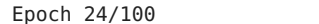






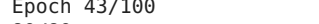






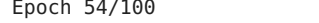


        # Define the learning rate scheduler (exponential increase)
        lr_scheduler = tf.keras.callbacks.LearningRateScheduler(
            lambda epoch: 1e-8 * 10**(epoch / 20) # Exponential growth
        )

        # Compile the model with the optimizer and loss function
        model.compile(
            optimizer=optimizer,
            loss=loss
        )

        # Fit the model with the learning rate scheduler
        history = model.fit(train_dataset, epochs=epochs, callbacks=[lr_scheduler])
        return history

# Example usage
tuned_history = tuning_learn_rate(
    model, X_train, y_train, optimizer=tf.keras.optimizers.Adam(), loss='binary_crossentropy'
)
```

```
Epoch 1/100
80/80 ————— 6s 13ms/step - loss: 1.0233 - learning_rate: 1.0000e-08
Epoch 2/100
80/80 ————— 1s 13ms/step - loss: 1.0207 - learning_rate: 1.1220e-08
Epoch 3/100
80/80 ————— 1s 11ms/step - loss: 1.0322 - learning_rate: 1.2589e-08
Epoch 4/100
80/80 ————— 1s 10ms/step - loss: 1.0249 - learning_rate: 1.4125e-08
Epoch 5/100
80/80 ————— 1s 11ms/step - loss: 1.0232 - learning_rate: 1.5849e-08
Epoch 6/100
80/80 ————— 1s 11ms/step - loss: 1.0188 - learning_rate: 1.7783e-08
Epoch 7/100
80/80 ————— 1s 10ms/step - loss: 1.0324 - learning_rate: 1.9953e-08
Epoch 8/100
80/80 ————— 1s 11ms/step - loss: 1.0155 - learning_rate: 2.2387e-08
Epoch 9/100
80/80 ————— 1s 10ms/step - loss: 1.0331 - learning_rate: 2.5119e-08
Epoch 10/100
80/80 ————— 1s 10ms/step - loss: 1.0210 - learning_rate: 2.8184e-08
Epoch 11/100
80/80 ————— 1s 10ms/step - loss: 1.0219 - learning_rate: 3.1623e-08
Epoch 12/100
80/80 ————— 2s 14ms/step - loss: 1.0172 - learning_rate: 3.5481e-08
Epoch 13/100
80/80 ————— 1s 14ms/step - loss: 1.0205 - learning_rate: 3.9811e-08
Epoch 14/100
80/80 ————— 1s 15ms/step - loss: 1.0282 - learning_rate: 4.4668e-08
Epoch 15/100
80/80 ————— 1s 15ms/step - loss: 1.0297 - learning_rate: 5.0119e-08
Epoch 16/100
80/80 ————— 1s 8ms/step - loss: 1.0256 - learning_rate: 5.6234e-08
```


Epoch 17/100
80/80  1s 11ms/step - loss: 1.0223 - learning_rate: 6.3096e-08
Epoch 18/100
80/80  1s 11ms/step - loss: 1.0132 - learning_rate: 7.0795e-08
Epoch 19/100
80/80  1s 5ms/step - loss: 1.0321 - learning_rate: 7.9433e-08
Epoch 20/100
80/80  1s 6ms/step - loss: 1.0366 - learning_rate: 8.9125e-08
Epoch 21/100
80/80  1s 5ms/step - loss: 1.0214 - learning_rate: 1.0000e-07
Epoch 22/100
80/80  0s 6ms/step - loss: 1.0249 - learning_rate: 1.1220e-07
Epoch 23/100
80/80  0s 5ms/step - loss: 1.0139 - learning_rate: 1.2589e-07
Epoch 24/100
80/80  1s 5ms/step - loss: 1.0235 - learning_rate: 1.4125e-07
Epoch 25/100
80/80  1s 6ms/step - loss: 1.0243 - learning_rate: 1.5849e-07
Epoch 26/100
80/80  0s 5ms/step - loss: 1.0291 - learning_rate: 1.7783e-07
Epoch 27/100
80/80  1s 6ms/step - loss: 1.0254 - learning_rate: 1.9953e-07
Epoch 28/100
80/80  1s 5ms/step - loss: 1.0145 - learning_rate: 2.2387e-07
Epoch 29/100
80/80  1s 5ms/step - loss: 1.0148 - learning_rate: 2.5119e-07
Epoch 30/100
80/80  1s 8ms/step - loss: 1.0196 - learning_rate: 2.8184e-07
Epoch 31/100
80/80  1s 10ms/step - loss: 1.0243 - learning_rate: 3.1623e-07
Epoch 32/100
80/80  1s 10ms/step - loss: 1.0250 - learning_rate: 3.5481e-07
Epoch 33/100
80/80  1s 5ms/step - loss: 1.0286 - learning_rate: 3.9811e-07
Epoch 34/100
80/80  1s 6ms/step - loss: 1.0345 - learning_rate: 4.4668e-07
Epoch 35/100
80/80  1s 6ms/step - loss: 1.0233 - learning_rate: 5.0119e-07
Epoch 36/100
80/80  1s 5ms/step - loss: 1.0232 - learning_rate: 5.6234e-07
Epoch 37/100
80/80  1s 5ms/step - loss: 1.0181 - learning_rate: 6.3096e-07
Epoch 38/100
80/80  0s 6ms/step - loss: 1.0299 - learning_rate: 7.0795e-07
Epoch 39/100
80/80  0s 6ms/step - loss: 1.0230 - learning_rate: 7.9433e-07
Epoch 40/100
80/80  0s 6ms/step - loss: 1.0220 - learning_rate: 8.9125e-07
Epoch 41/100
80/80  0s 6ms/step - loss: 1.0297 - learning_rate: 1.0000e-06
Epoch 42/100
80/80  1s 6ms/step - loss: 1.0267 - learning_rate: 1.1220e-06
Epoch 43/100
80/80  1s 6ms/step - loss: 1.0297 - learning_rate: 1.2589e-06
Epoch 44/100
80/80  1s 5ms/step - loss: 1.0242 - learning_rate: 1.4125e-06
Epoch 45/100
80/80  0s 6ms/step - loss: 1.0324 - learning_rate: 1.5849e-06
Epoch 46/100
80/80  1s 5ms/step - loss: 1.0323 - learning_rate: 1.7783e-06
Epoch 47/100
80/80  0s 6ms/step - loss: 1.0356 - learning_rate: 1.9953e-06
Epoch 48/100
80/80  0s 5ms/step - loss: 1.0256 - learning_rate: 2.2387e-06
Epoch 49/100
80/80  1s 6ms/step - loss: 1.0206 - learning_rate: 2.5119e-06
Epoch 50/100
80/80  1s 5ms/step - loss: 1.0333 - learning_rate: 2.8184e-06
Epoch 51/100
80/80  1s 10ms/step - loss: 1.0158 - learning_rate: 3.1623e-06
Epoch 52/100
80/80  1s 10ms/step - loss: 1.0163 - learning_rate: 3.5481e-06
Epoch 53/100
80/80  1s 8ms/step - loss: 1.0304 - learning_rate: 3.9811e-06
Epoch 54/100
80/80  0s 5ms/step - loss: 1.0158 - learning_rate: 4.4668e-06
Epoch 55/100
80/80  1s 6ms/step - loss: 1.0039 - learning_rate: 5.0119e-06
Epoch 56/100
80/80  1s 5ms/step - loss: 1.0127 - learning_rate: 5.6234e-06
Epoch 57/100
80/80  1s 5ms/step - loss: 0.9992 - learning_rate: 6.3096e-06
Epoch 58/100
80/80  1s 6ms/step - loss: 0.9882 - learning_rate: 7.0795e-06
Epoch 59/100
80/80  0s 5ms/step - loss: 0.9902 - learning_rate: 7.9433e-06
Epoch 60/100
80/80  1s 6ms/step - loss: 0.9839 - learning_rate: 8.9125e-06
Epoch 61/100

```

80/80 ————— 0s 5ms/step - loss: 0.9709 - learning_rate: 1.0000e-05
Epoch 62/100
80/80 ————— 1s 6ms/step - loss: 0.9464 - learning_rate: 1.1220e-05
Epoch 63/100
80/80 ————— 1s 5ms/step - loss: 0.9408 - learning_rate: 1.2589e-05
Epoch 64/100
80/80 ————— 1s 5ms/step - loss: 0.9237 - learning_rate: 1.4125e-05
Epoch 65/100
80/80 ————— 1s 6ms/step - loss: 0.9007 - learning_rate: 1.5849e-05
Epoch 66/100
80/80 ————— 1s 6ms/step - loss: 0.9043 - learning_rate: 1.7783e-05
Epoch 67/100
80/80 ————— 1s 6ms/step - loss: 0.8718 - learning_rate: 1.9953e-05
Epoch 68/100
80/80 ————— 1s 6ms/step - loss: 0.8349 - learning_rate: 2.2387e-05
Epoch 69/100
80/80 ————— 1s 6ms/step - loss: 0.8053 - learning_rate: 2.5119e-05
Epoch 70/100
80/80 ————— 1s 8ms/step - loss: 0.7787 - learning_rate: 2.8184e-05
Epoch 71/100
80/80 ————— 1s 10ms/step - loss: 0.7432 - learning_rate: 3.1623e-05
Epoch 72/100
80/80 ————— 1s 10ms/step - loss: 0.7056 - learning_rate: 3.5481e-05
Epoch 73/100
80/80 ————— 1s 11ms/step - loss: 0.6511 - learning_rate: 3.9811e-05
Epoch 74/100
80/80 ————— 1s 6ms/step - loss: 0.6035 - learning_rate: 4.4668e-05
Epoch 75/100
80/80 ————— 1s 5ms/step - loss: 0.5657 - learning_rate: 5.0119e-05
Epoch 76/100
80/80 ————— 0s 6ms/step - loss: 0.5150 - learning_rate: 5.6234e-05
Epoch 77/100
80/80 ————— 1s 6ms/step - loss: 0.4717 - learning_rate: 6.3096e-05
Epoch 78/100
80/80 ————— 0s 5ms/step - loss: 0.4254 - learning_rate: 7.0795e-05
Epoch 79/100
80/80 ————— 0s 5ms/step - loss: 0.4001 - learning_rate: 7.9433e-05
Epoch 80/100
80/80 ————— 1s 6ms/step - loss: 0.3654 - learning_rate: 8.9125e-05
Epoch 81/100
80/80 ————— 0s 6ms/step - loss: 0.3415 - learning_rate: 1.0000e-04
Epoch 82/100
80/80 ————— 0s 6ms/step - loss: 0.3244 - learning_rate: 1.1220e-04
Epoch 83/100
80/80 ————— 1s 6ms/step - loss: 0.3043 - learning_rate: 1.2589e-04
Epoch 84/100
80/80 ————— 1s 5ms/step - loss: 0.2898 - learning_rate: 1.4125e-04
Epoch 85/100
80/80 ————— 1s 6ms/step - loss: 0.2785 - learning_rate: 1.5849e-04
Epoch 86/100
80/80 ————— 1s 6ms/step - loss: 0.2744 - learning_rate: 1.7783e-04
Epoch 87/100
80/80 ————— 0s 6ms/step - loss: 0.2557 - learning_rate: 1.9953e-04
Epoch 88/100
80/80 ————— 1s 6ms/step - loss: 0.2479 - learning_rate: 2.2387e-04
Epoch 89/100
80/80 ————— 0s 5ms/step - loss: 0.2384 - learning_rate: 2.5119e-04
Epoch 90/100
80/80 ————— 0s 6ms/step - loss: 0.2238 - learning_rate: 2.8184e-04
Epoch 91/100
80/80 ————— 1s 5ms/step - loss: 0.2172 - learning_rate: 3.1623e-04
Epoch 92/100
80/80 ————— 1s 10ms/step - loss: 0.2081 - learning_rate: 3.5481e-04
Epoch 93/100
80/80 ————— 1s 10ms/step - loss: 0.1982 - learning_rate: 3.9811e-04
Epoch 94/100
80/80 ————— 1s 11ms/step - loss: 0.1848 - learning_rate: 4.4668e-04
Epoch 95/100
80/80 ————— 1s 11ms/step - loss: 0.1725 - learning_rate: 5.0119e-04
Epoch 96/100
80/80 ————— 0s 6ms/step - loss: 0.1585 - learning_rate: 5.6234e-04
Epoch 97/100
80/80 ————— 0s 6ms/step - loss: 0.1497 - learning_rate: 6.3096e-04
Epoch 98/100
80/80 ————— 1s 6ms/step - loss: 0.1313 - learning_rate: 7.0795e-04
Epoch 99/100
80/80 ————— 0s 5ms/step - loss: 0.1191 - learning_rate: 7.9433e-04
Epoch 100/100
80/80 ————— 0s 6ms/step - loss: 0.1082 - learning_rate: 8.9125e-04

```

```
In [81]: plot_learning_rate(tuned_history)
```



```
In [82]: # Compiling the Network

model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=2e-4), loss='binary_crossentropy', metrics=['acc'])
# learning_rate=4e-4
```

```
In [83]: model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 256)	7,936
batch_normalization_2 (BatchNormalization)	(None, 256)	1,024
dropout_4 (Dropout)	(None, 256)	0
dense_9 (Dense)	(None, 256)	65,792
dropout_5 (Dropout)	(None, 256)	0
dense_10 (Dense)	(None, 128)	32,896
dense_11 (Dense)	(None, 1)	129

Total params: 107,777 (421.00 KB)

Trainable params: 107,265 (419.00 KB)

Non-trainable params: 512 (2.00 KB)

```
In [84]: # Training the data on the network

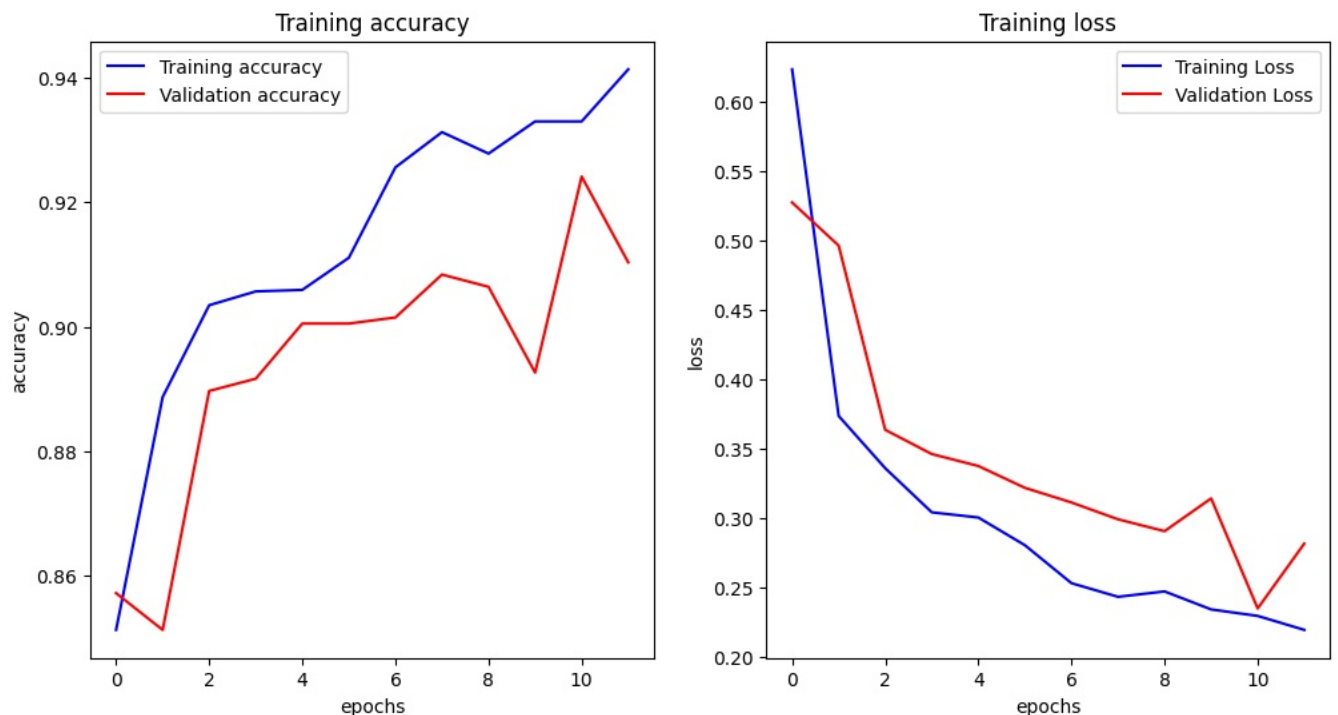
history = model.fit(x=X_train, y=y_train, validation_split=0.2, batch_size=32, epochs=50, callbacks=[EarlyStopping])
```

```

Epoch 1/50
127/127 — 3s 7ms/step - accuracy: 0.7858 - loss: 0.9670 - val_accuracy: 0.8571 - val_loss: 0.5272
Epoch 2/50
127/127 — 1s 5ms/step - accuracy: 0.8886 - loss: 0.4006 - val_accuracy: 0.8512 - val_loss: 0.4961
Epoch 3/50
127/127 — 1s 5ms/step - accuracy: 0.8983 - loss: 0.3592 - val_accuracy: 0.8897 - val_loss: 0.3633
Epoch 4/50
127/127 — 2s 8ms/step - accuracy: 0.9018 - loss: 0.3097 - val_accuracy: 0.8916 - val_loss: 0.3459
Epoch 5/50
127/127 — 1s 8ms/step - accuracy: 0.9086 - loss: 0.2991 - val_accuracy: 0.9005 - val_loss: 0.3372
Epoch 6/50
127/127 — 1s 9ms/step - accuracy: 0.9114 - loss: 0.2835 - val_accuracy: 0.9005 - val_loss: 0.3215
Epoch 7/50
127/127 — 1s 9ms/step - accuracy: 0.9239 - loss: 0.2501 - val_accuracy: 0.9015 - val_loss: 0.3110
Epoch 8/50
127/127 — 1s 6ms/step - accuracy: 0.9322 - loss: 0.2507 - val_accuracy: 0.9084 - val_loss: 0.2988
Epoch 9/50
127/127 — 1s 5ms/step - accuracy: 0.9279 - loss: 0.2456 - val_accuracy: 0.9064 - val_loss: 0.2903
Epoch 10/50
127/127 — 1s 5ms/step - accuracy: 0.9327 - loss: 0.2319 - val_accuracy: 0.8926 - val_loss: 0.3139
Epoch 11/50
127/127 — 1s 5ms/step - accuracy: 0.9371 - loss: 0.2196 - val_accuracy: 0.9241 - val_loss: 0.2346
Epoch 12/50
124/127 — 0s 4ms/step - accuracy: 0.9379 - loss: 0.2258
Reached 94% accuracy so cancelling training!
127/127 — 1s 5ms/step - accuracy: 0.9380 - loss: 0.2256 - val_accuracy: 0.9103 - val_loss: 0.2813

```

In [85]: `plot_loss_acc(history)`



In [86]: `y_pred = model.predict(X_test)`
`y_pred = (y_pred > 0.5)`
`print(np.concatenate((y_pred.reshape(-1, 1), y_test.reshape(-1, 1)), 1))`

```

44/44 — 0s 3ms/step
[[0 0]
 [0 0]
 [1 1]
 ...
 [0 0]
 [0 0]
 [0 0]]

```

In [87]: `from sklearn.metrics import confusion_matrix, accuracy_score`
`cm = confusion_matrix(y_test, y_pred)`
`print(cm)`
`print(accuracy_score(y_test, y_pred))`

```
[[735 298]
 [ 96 278]]
0.7199715707178393
```

```
In [88]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.71	0.79	1033
1	0.48	0.74	0.59	374
accuracy			0.72	1407
macro avg	0.68	0.73	0.69	1407
weighted avg	0.78	0.72	0.73	1407