

nf9ri6zin

January 4, 2025

## 0.1 Importing Libraries

```
[ ]: !pip install catboost
```

```
[ ]: !pip uninstall -y scikit-learn  
!pip install scikit-learn==1.3.1
```

```
[ ]: !pip install imblearn
```

```
[4]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.preprocessing import LabelEncoder  
from sklearn.preprocessing import OneHotEncoder  
from sklearn.compose import ColumnTransformer  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
from sklearn.metrics import accuracy_score, confusion_matrix  
from xgboost import XGBClassifier  
from sklearn.svm import SVC  
from catboost import CatBoostClassifier  
from sklearn.model_selection import RandomizedSearchCV  
from sklearn.model_selection import cross_val_score
```

## 0.2 Reading The Dataset

```
[5]: # Loading the dataset  
file_path = './WA_Fn-UseC_-Telco-Customer-Churn.csv'  
  
data = pd.read_csv(file_path)  
print(f'dataset contains {data.shape[0]} rows and {data.shape[1]} columns')
```

dataset contains 7043 rows and 21 columns

```
[6]: data.head(10)
```

```
[6]:  customerID  gender  Senior_Citizen  Is_Married  Dependents  tenure  \
0  7590-VHVEG  Female                0           Yes           No           1
1  5575-GNVDE   Male                0           No            No           34
2  3668-QPYBK   Male                0           No            No           2
3  7795-CFOCW   Male                0           No            No           45
4  9237-HQITU   Female               0           No            No           2
5  9305-CDSKC   Female               0           No            No           8
6  1452-KIOVK   Male                0           No            Yes          22
7  6713-OKOMC   Female               0           No            No           10
8  7892-POOKP   Female               0           Yes           No           28
9  6388-TABGU   Male                0           No            Yes          62
```

```
Phone_Service      Dual  Internet_Service  Online_Security  ...  \
0           No  No phone service           DSL           No  ...
1           Yes           No           DSL           Yes  ...
2           Yes           No           DSL           Yes  ...
3           No  No phone service           DSL           Yes  ...
4           Yes           No  Fiber optic           No  ...
5           Yes           Yes  Fiber optic           No  ...
6           Yes           Yes  Fiber optic           No  ...
7           No  No phone service           DSL           Yes  ...
8           Yes           Yes  Fiber optic           No  ...
9           Yes           No           DSL           Yes  ...
```

```
Device_Protection  Tech_Support  Streaming_TV  Streaming_Movies  \
0           No           No           No           No
1           Yes           No           No           No
2           No           No           No           No
3           Yes           Yes           No           No
4           No           No           No           No
5           Yes           No           Yes           Yes
6           No           No           Yes           No
7           No           No           No           No
8           Yes           Yes           Yes           Yes
9           No           No           No           No
```

```
Contract  Paperless_Billing      Payment_Method  \
0  Month-to-month           Yes  Electronic check
1    One year              No    Mailed check
2  Month-to-month           Yes    Mailed check
3    One year              No  Bank transfer (automatic)
4  Month-to-month           Yes    Electronic check
5  Month-to-month           Yes    Electronic check
6  Month-to-month           Yes  Credit card (automatic)
7  Month-to-month           No    Mailed check
8  Month-to-month           Yes    Electronic check
9    One year              No  Bank transfer (automatic)
```

|   | Monthly_Charges | Total_Charges | Churn |
|---|-----------------|---------------|-------|
| 0 | 29.85           | 29.85         | No    |
| 1 | 56.95           | 1889.5        | No    |
| 2 | 53.85           | 108.15        | Yes   |
| 3 | 42.30           | 1840.75       | No    |
| 4 | 70.70           | 151.65        | Yes   |
| 5 | 99.65           | 820.5         | Yes   |
| 6 | 89.10           | 1949.4        | No    |
| 7 | 29.75           | 301.9         | No    |
| 8 | 104.80          | 3046.05       | Yes   |
| 9 | 56.15           | 3487.95       | No    |

[10 rows x 21 columns]

```
[7]: data.tail()
```

```
[7]:      customerID  gender  Senior_Citizen  Is_Married  Dependents  tenure  \
7038  6840-RESVB    Male                0         Yes         Yes       24
7039  2234-XADUH   Female                0         Yes         Yes       72
7040  4801-JZAZL   Female                0         Yes         Yes       11
7041  8361-LTMKD    Male                1         Yes         No        4
7042  3186-AJIEK    Male                0         No          No       66
```

|      | Phone_Service | Dual             | Internet_Service | Online_Security | ... | \ |
|------|---------------|------------------|------------------|-----------------|-----|---|
| 7038 | Yes           | Yes              | DSL              | Yes             | ... |   |
| 7039 | Yes           | Yes              | Fiber optic      | No              | ... |   |
| 7040 | No            | No phone service | DSL              | Yes             | ... |   |
| 7041 | Yes           | Yes              | Fiber optic      | No              | ... |   |
| 7042 | Yes           | No               | Fiber optic      | Yes             | ... |   |

|      | Device_Protection | Tech_Support | Streaming_TV | Streaming_Movies | \ |
|------|-------------------|--------------|--------------|------------------|---|
| 7038 | Yes               | Yes          | Yes          | Yes              |   |
| 7039 | Yes               | No           | Yes          | Yes              |   |
| 7040 | No                | No           | No           | No               |   |
| 7041 | No                | No           | No           | No               |   |
| 7042 | Yes               | Yes          | Yes          | Yes              |   |

|      | Contract       | Paperless_Billing | Payment_Method            | \ |
|------|----------------|-------------------|---------------------------|---|
| 7038 | One year       | Yes               | Mailed check              |   |
| 7039 | One year       | Yes               | Credit card (automatic)   |   |
| 7040 | Month-to-month | Yes               | Electronic check          |   |
| 7041 | Month-to-month | Yes               | Mailed check              |   |
| 7042 | Two year       | Yes               | Bank transfer (automatic) |   |

|      | Monthly_Charges | Total_Charges | Churn |
|------|-----------------|---------------|-------|
| 7038 | 84.80           | 1990.5        | No    |

|      |        |        |     |
|------|--------|--------|-----|
| 7039 | 103.20 | 7362.9 | No  |
| 7040 | 29.60  | 346.45 | No  |
| 7041 | 74.40  | 306.6  | Yes |
| 7042 | 105.65 | 6844.5 | No  |

[5 rows x 21 columns]

### 0.3 Get Information about the dataset:

```
[8]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   Senior_Citizen        7043 non-null   int64
3   Is_Married            7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                 7043 non-null   int64
6   Phone_Service         7043 non-null   object
7   Dual                  7043 non-null   object
8   Internet_Service      7043 non-null   object
9   Online_Security       7043 non-null   object
10  Online_Backup         7043 non-null   object
11  Device_Protection     7043 non-null   object
12  Tech_Support          7043 non-null   object
13  Streaming_TV          7043 non-null   object
14  Streaming_Movies      7043 non-null   object
15  Contract              7043 non-null   object
16  Paperless_Billing     7043 non-null   object
17  Payment_Method        7043 non-null   object
18  Monthly_Charges       7043 non-null   float64
19  Total_Charges         7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

### 0.4 Covertng Some of the Object Datatype Columns to Numerical

```
[9]: # Total_charges column should be of numerical type
data['Total_Charges'] = data['Total_Charges'].apply(pd.to_numeric,
↳errors='coerce') # invalid parsing will be set as NaN
```

## 0.5 Checking Null Values

```
[10]: data.isnull().sum()
```

```
[10]: customerID      0
      gender          0
      Senior_Citizen  0
      Is_Married      0
      Dependents      0
      tenure          0
      Phone_Service   0
      Dual            0
      Internet_Service 0
      Online_Security 0
      Online_Backup   0
      Device_Protection 0
      Tech_Support    0
      Streaming_TV     0
      Streaming_Movies 0
      Contract        0
      Paperless_Billing 0
      Payment_Method   0
      Monthly_Charges  0
      Total_Charges    11
      Churn            0
      dtype: int64
```

```
[11]: # Check rows where Total_Charges is NaN
      nan_rows = data[data['Total_Charges'].isna()]
      print(nan_rows)
```

|      | customerID | gender | Senior_Citizen | Is_Married | Dependents | tenure | \ |
|------|------------|--------|----------------|------------|------------|--------|---|
| 488  | 4472-LVYGI | Female | 0              | Yes        | Yes        | 0      |   |
| 753  | 3115-CZMZD | Male   | 0              | No         | Yes        | 0      |   |
| 936  | 5709-LVOEQ | Female | 0              | Yes        | Yes        | 0      |   |
| 1082 | 4367-NUYAO | Male   | 0              | Yes        | Yes        | 0      |   |
| 1340 | 1371-DWPAZ | Female | 0              | Yes        | Yes        | 0      |   |
| 3331 | 7644-OMVMY | Male   | 0              | Yes        | Yes        | 0      |   |
| 3826 | 3213-VVOLG | Male   | 0              | Yes        | Yes        | 0      |   |
| 4380 | 2520-SGTTA | Female | 0              | Yes        | Yes        | 0      |   |
| 5218 | 2923-ARZLG | Male   | 0              | Yes        | Yes        | 0      |   |
| 6670 | 4075-WKNIU | Female | 0              | Yes        | Yes        | 0      |   |
| 6754 | 2775-SEFEE | Male   | 0              | No         | Yes        | 0      |   |

|     | Phone_Service | Dual             | Internet_Service | Online_Security     | \ |
|-----|---------------|------------------|------------------|---------------------|---|
| 488 | No            | No phone service | DSL              | Yes                 |   |
| 753 | Yes           | No               | No               | No internet service |   |
| 936 | Yes           | No               | DSL              | Yes                 |   |

|      |     |                  |     |                     |
|------|-----|------------------|-----|---------------------|
| 1082 | Yes | Yes              | No  | No internet service |
| 1340 | No  | No phone service | DSL | Yes                 |
| 3331 | Yes | No               | No  | No internet service |
| 3826 | Yes | Yes              | No  | No internet service |
| 4380 | Yes | No               | No  | No internet service |
| 5218 | Yes | No               | No  | No internet service |
| 6670 | Yes | Yes              | DSL | No                  |
| 6754 | Yes | Yes              | DSL | Yes                 |

|      |     |                     |                     |                     |                     |              |   |
|------|-----|---------------------|---------------------|---------------------|---------------------|--------------|---|
|      | ... | Device_Protection   |                     | Tech_Support        |                     | Streaming_TV | \ |
| 488  | ... | Yes                 |                     | Yes                 |                     | Yes          |   |
| 753  | ... | No internet service | No internet service | No internet service | No internet service |              |   |
| 936  | ... | Yes                 |                     | No                  |                     | Yes          |   |
| 1082 | ... | No internet service | No internet service | No internet service | No internet service |              |   |
| 1340 | ... | Yes                 |                     | Yes                 |                     | Yes          |   |
| 3331 | ... | No internet service | No internet service | No internet service | No internet service |              |   |
| 3826 | ... | No internet service | No internet service | No internet service | No internet service |              |   |
| 4380 | ... | No internet service | No internet service | No internet service | No internet service |              |   |
| 5218 | ... | No internet service | No internet service | No internet service | No internet service |              |   |
| 6670 | ... | Yes                 |                     | Yes                 |                     | Yes          |   |
| 6754 | ... | No                  |                     | Yes                 |                     | No           |   |

|      |                     |                  |          |                   |   |
|------|---------------------|------------------|----------|-------------------|---|
|      |                     | Streaming_Movies | Contract | Paperless_Billing | \ |
| 488  |                     | No               | Two year | Yes               |   |
| 753  | No internet service |                  | Two year | No                |   |
| 936  |                     | Yes              | Two year | No                |   |
| 1082 | No internet service |                  | Two year | No                |   |
| 1340 |                     | No               | Two year | No                |   |
| 3331 | No internet service |                  | Two year | No                |   |
| 3826 | No internet service |                  | Two year | No                |   |
| 4380 | No internet service |                  | Two year | No                |   |
| 5218 | No internet service |                  | One year | Yes               |   |
| 6670 |                     | No               | Two year | No                |   |
| 6754 |                     | No               | Two year | Yes               |   |

|      |                           |                |                 |               |       |
|------|---------------------------|----------------|-----------------|---------------|-------|
|      |                           | Payment_Method | Monthly_Charges | Total_Charges | Churn |
| 488  | Bank transfer (automatic) |                | 52.55           | NaN           | No    |
| 753  | Mailed check              |                | 20.25           | NaN           | No    |
| 936  | Mailed check              |                | 80.85           | NaN           | No    |
| 1082 | Mailed check              |                | 25.75           | NaN           | No    |
| 1340 | Credit card (automatic)   |                | 56.05           | NaN           | No    |
| 3331 | Mailed check              |                | 19.85           | NaN           | No    |
| 3826 | Mailed check              |                | 25.35           | NaN           | No    |
| 4380 | Mailed check              |                | 20.00           | NaN           | No    |
| 5218 | Mailed check              |                | 19.70           | NaN           | No    |
| 6670 | Mailed check              |                | 73.35           | NaN           | No    |
| 6754 | Bank transfer (automatic) |                | 61.90           | NaN           | No    |

[11 rows x 21 columns]

```
[12]: # Due to the dataset being big i decided to drop the rows

# Drop rows with NaN in the 'Total_Charges' column
data = data.dropna(subset=['Total_Charges'])
```

```
[13]: data['Total_Charges'].isnull().sum()
```

[13]: 0

## 0.6 Checking Duplicates

```
[14]: data_dup = data.duplicated().any()
print(data_dup)
# data.drop_duplicates()
```

False

## 0.7 Overall statistics about the dataset

```
[15]: data.describe()
```

```
[15]:
```

|       | Senior_Citizen | tenure      | Monthly_Charges | Total_Charges |
|-------|----------------|-------------|-----------------|---------------|
| count | 7032.000000    | 7032.000000 | 7032.000000     | 7032.000000   |
| mean  | 0.162400       | 32.421786   | 64.798208       | 2283.300441   |
| std   | 0.368844       | 24.545260   | 30.085974       | 2266.771362   |
| min   | 0.000000       | 1.000000    | 18.250000       | 18.800000     |
| 25%   | 0.000000       | 9.000000    | 35.587500       | 401.450000    |
| 50%   | 0.000000       | 29.000000   | 70.350000       | 1397.475000   |
| 75%   | 0.000000       | 55.000000   | 89.862500       | 3794.737500   |
| max   | 1.000000       | 72.000000   | 118.750000      | 8684.800000   |

## 0.8 Splitting the Dataset into Features and a Target

```
[16]: X = data.iloc[:, 1:-1] # Excluding the customerid and the churn
y = data.iloc[:, -1] # Churn
```

```
[17]: print(X.shape)
print(y.shape)
```

(7032, 19)

(7032,)

```
[18]: X.head(5)
```

```
[18]:  gender  Senior_Citizen  Is_Married  Dependents  tenure  Phone_Service  \
0  Female                0          Yes          No         1           No
1   Male                0          No          No        34          Yes
2   Male                0          No          No         2          Yes
3   Male                0          No          No        45           No
4  Female                0          No          No         2          Yes
```

```

          Dual  Internet_Service  Online_Security  Online_Backup  \
0  No phone service              DSL              No              Yes
1                No              DSL              Yes              No
2                No              DSL              Yes              Yes
3  No phone service              DSL              Yes              No
4                No      Fiber optic              No              No
```

```

Device_Protection  Tech_Support  Streaming_TV  Streaming_Movies  \
0                No            No            No            No
1                Yes            No            No            No
2                No            No            No            No
3                Yes            Yes            No            No
4                No            No            No            No
```

```

          Contract  Paperless_Billing              Payment_Method  \
0  Month-to-month              Yes          Electronic check
1    One year                No          Mailed check
2  Month-to-month              Yes          Mailed check
3    One year                No  Bank transfer (automatic)
4  Month-to-month              Yes          Electronic check
```

```

Monthly_Charges  Total_Charges
0          29.85          29.85
1          56.95        1889.50
2          53.85         108.15
3          42.30        1840.75
4          70.70         151.65
```

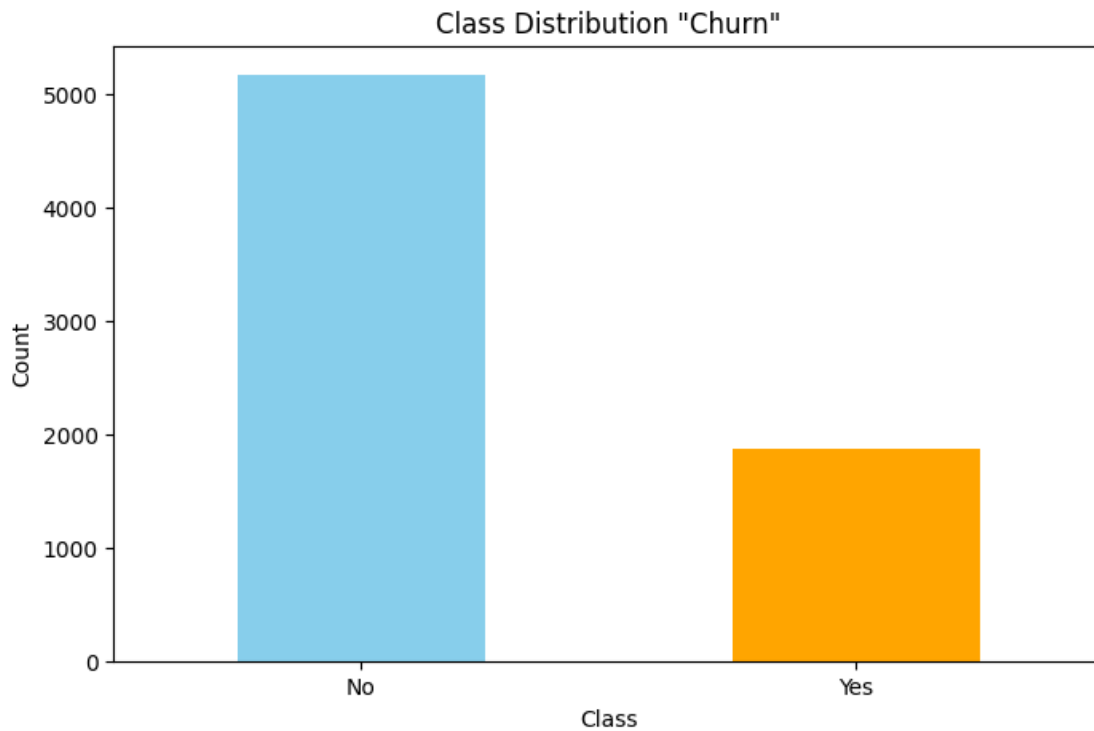
```
[19]: y.head(5)
```

```
[19]: 0    No
1    No
2    Yes
3    No
4    Yes
Name: Churn, dtype: object
```

```
[20]: # Get the count of each class
class_counts = y.value_counts() # Counts of unique values
```




```
# Plot the counts as a bar chart
plt.figure(figsize=(8, 5))
class_counts.plot(kind='bar', color=['skyblue', 'orange'])
plt.title('Class Distribution "Churn"')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.show()
```



```
[21]: class_counts.plot.pie(autopct='%1.2f%%', colors=['skyblue', 'orange'])
plt.title('Class Distribution "Churn"')
plt.ylabel('')
plt.show()
# Looks Like an Imbalanced Dataset
```

Class Distribution "Churn"



## 0.9 Encoding Categorical Data

```
[22]: # Checking unique values to choose which technique to apply
      should_be_one_hot_encoded = []
      should_be_label_encoded = []

      for col in X.columns:
          if X[col].dtypes == 'object': # Exclude numerical values
              print(f'{col}: {X[col].unique()}')
              if len(X[col].unique()) > 2:
                  should_be_one_hot_encoded.append(col)
              else:
                  should_be_label_encoded.append(col)

      print('\nOne-Hot Encoded : ', should_be_one_hot_encoded, '\n')
      print('Label Encoded : ', should_be_label_encoded)
```

```
gender: ['Female' 'Male']
Is_Married: ['Yes' 'No']
Dependents: ['No' 'Yes']
Phone_Service: ['No' 'Yes']
Dual: ['No phone service' 'No' 'Yes']
```

```

Internet_Service: ['DSL' 'Fiber optic' 'No']
Online_Security: ['No' 'Yes' 'No internet service']
Online_Backup: ['Yes' 'No' 'No internet service']
Device_Protection: ['No' 'Yes' 'No internet service']
Tech_Support: ['No' 'Yes' 'No internet service']
Streaming_TV: ['No' 'Yes' 'No internet service']
Streaming_Movies: ['No' 'Yes' 'No internet service']
Contract: ['Month-to-month' 'One year' 'Two year']
Paperless_Billing: ['Yes' 'No']
Payment_Method: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']

```

```

One-Hot Encoded : ['Dual', 'Internet_Service', 'Online_Security',
'Online_Backup', 'Device_Protection', 'Tech_Support', 'Streaming_TV',
'Streaming_Movies', 'Contract', 'Payment_Method']

```

```

Label Encoded : ['gender', 'Is_Married', 'Dependents', 'Phone_Service',
'Paperless_Billing']

```

### 0.9.1 Apply Label Encoding

```

[23]: le = LabelEncoder()
      for col in should_be_label_encoded:
          X[col] = le.fit_transform(X[col]) # Apply label encoding for each column

      for col in should_be_label_encoded:
          print(f'{col}: {X[col].unique()}')

```

```

gender: [0 1]
Is_Married: [1 0]
Dependents: [0 1]
Phone_Service: [0 1]
Paperless_Billing: [1 0]

```

```

[24]: # Label Encoding the Target
      y = le.fit_transform(y)
      print(y)

```

```

[0 0 1 ... 0 1 0]

```

### 0.9.2 Apply One-Hot Encoding

```

[25]: # Get the indexes of columns to transform
      hot_encode_indexes = X.columns.get_indexer(should_be_one_hot_encoded)
      print(hot_encode_indexes)
      ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(),
      ↪hot_encode_indexes)], remainder='passthrough')
      # Fit and transform the data

```

```
X_transformed = np.array(ct.fit_transform(X))
print(X_transformed)
```

```
[ 6  7  8  9 10 11 12 13 14 16]
[[0.0000e+00 1.0000e+00 0.0000e+00 ... 1.0000e+00 2.9850e+01 2.9850e+01]
 [1.0000e+00 0.0000e+00 0.0000e+00 ... 0.0000e+00 5.6950e+01 1.8895e+03]
 [1.0000e+00 0.0000e+00 0.0000e+00 ... 1.0000e+00 5.3850e+01 1.0815e+02]
 ...
 [0.0000e+00 1.0000e+00 0.0000e+00 ... 1.0000e+00 2.9600e+01 3.4645e+02]
 [0.0000e+00 0.0000e+00 1.0000e+00 ... 1.0000e+00 7.4400e+01 3.0660e+02]
 [1.0000e+00 0.0000e+00 0.0000e+00 ... 1.0000e+00 1.0565e+02 6.8445e+03]]
```

```
[26]: # Get the feature names for the one-hot encoded columns
encoder = ct.transformers_[0][1] # The encoder used for one-hot encoding
encoded_feature_names = encoder.
    ↳ get_feature_names_out(input_features=should_be_one_hot_encoded)

# Create a DataFrame with the transformed data
# Concatenate the new one-hot encoded feature names and original columns that
    ↳ weren't transformed
X_transformed_df = pd.DataFrame(X_transformed, columns=np.
    ↳ concatenate([encoded_feature_names, X.columns.
    ↳ difference(should_be_one_hot_encoded)]))

# Show the resulting DataFrame
print(X_transformed_df)
```

|      | Dual_No | Dual_No phone service | Dual_Yes | Internet_Service_DSL | \ |
|------|---------|-----------------------|----------|----------------------|---|
| 0    | 0.0     | 1.0                   | 0.0      | 1.0                  |   |
| 1    | 1.0     | 0.0                   | 0.0      | 1.0                  |   |
| 2    | 1.0     | 0.0                   | 0.0      | 1.0                  |   |
| 3    | 0.0     | 1.0                   | 0.0      | 1.0                  |   |
| 4    | 1.0     | 0.0                   | 0.0      | 0.0                  |   |
| ...  | ...     | ...                   | ...      | ...                  |   |
| 7027 | 0.0     | 0.0                   | 1.0      | 1.0                  |   |
| 7028 | 0.0     | 0.0                   | 1.0      | 0.0                  |   |
| 7029 | 0.0     | 1.0                   | 0.0      | 1.0                  |   |
| 7030 | 0.0     | 0.0                   | 1.0      | 0.0                  |   |
| 7031 | 1.0     | 0.0                   | 0.0      | 0.0                  |   |

|     | Internet_Service_Fiber optic | Internet_Service_No | Online_Security_No | \ |
|-----|------------------------------|---------------------|--------------------|---|
| 0   | 0.0                          | 0.0                 | 1.0                |   |
| 1   | 0.0                          | 0.0                 | 0.0                |   |
| 2   | 0.0                          | 0.0                 | 0.0                |   |
| 3   | 0.0                          | 0.0                 | 0.0                |   |
| 4   | 1.0                          | 0.0                 | 1.0                |   |
| ... | ...                          | ...                 | ...                |   |

|      |     |     |     |
|------|-----|-----|-----|
| 7027 | 0.0 | 0.0 | 0.0 |
| 7028 | 1.0 | 0.0 | 1.0 |
| 7029 | 0.0 | 0.0 | 0.0 |
| 7030 | 1.0 | 0.0 | 1.0 |
| 7031 | 1.0 | 0.0 | 0.0 |

|      | Online_Security_No internet service | Online_Security_Yes \ |
|------|-------------------------------------|-----------------------|
| 0    | 0.0                                 | 0.0                   |
| 1    | 0.0                                 | 1.0                   |
| 2    | 0.0                                 | 1.0                   |
| 3    | 0.0                                 | 1.0                   |
| 4    | 0.0                                 | 0.0                   |
| ...  | ...                                 | ...                   |
| 7027 | 0.0                                 | 1.0                   |
| 7028 | 0.0                                 | 0.0                   |
| 7029 | 0.0                                 | 1.0                   |
| 7030 | 0.0                                 | 0.0                   |
| 7031 | 0.0                                 | 1.0                   |

|      | Online_Backup_No ... | Payment_Method_Mailed check | Dependents \ |
|------|----------------------|-----------------------------|--------------|
| 0    | 0.0 ...              | 0.0                         | 0.0          |
| 1    | 1.0 ...              | 1.0                         | 1.0          |
| 2    | 0.0 ...              | 1.0                         | 1.0          |
| 3    | 1.0 ...              | 0.0                         | 1.0          |
| 4    | 1.0 ...              | 0.0                         | 0.0          |
| ...  | ... ...              | ...                         | ...          |
| 7027 | 1.0 ...              | 1.0                         | 1.0          |
| 7028 | 0.0 ...              | 0.0                         | 0.0          |
| 7029 | 1.0 ...              | 0.0                         | 0.0          |
| 7030 | 1.0 ...              | 1.0                         | 1.0          |
| 7031 | 1.0 ...              | 0.0                         | 1.0          |

|      | Is_Married | Monthly_Charges | Paperless_Billing | Phone_Service \ |
|------|------------|-----------------|-------------------|-----------------|
| 0    | 0.0        | 1.0             | 0.0               | 1.0             |
| 1    | 0.0        | 0.0             | 0.0               | 34.0            |
| 2    | 0.0        | 0.0             | 0.0               | 2.0             |
| 3    | 0.0        | 0.0             | 0.0               | 45.0            |
| 4    | 0.0        | 0.0             | 0.0               | 2.0             |
| ...  | ...        | ...             | ...               | ...             |
| 7027 | 0.0        | 1.0             | 1.0               | 24.0            |
| 7028 | 0.0        | 1.0             | 1.0               | 72.0            |
| 7029 | 0.0        | 1.0             | 1.0               | 11.0            |
| 7030 | 1.0        | 1.0             | 0.0               | 4.0             |
| 7031 | 0.0        | 0.0             | 0.0               | 66.0            |

|   | Senior_Citizen | Total_Charges | gender | tenure  |
|---|----------------|---------------|--------|---------|
| 0 | 0.0            | 1.0           | 29.85  | 29.85   |
| 1 | 1.0            | 0.0           | 56.95  | 1889.50 |

|      |     |     |        |         |
|------|-----|-----|--------|---------|
| 2    | 1.0 | 1.0 | 53.85  | 108.15  |
| 3    | 0.0 | 0.0 | 42.30  | 1840.75 |
| 4    | 1.0 | 1.0 | 70.70  | 151.65  |
| ...  | ... | ... | ...    | ...     |
| 7027 | 1.0 | 1.0 | 84.80  | 1990.50 |
| 7028 | 1.0 | 1.0 | 103.20 | 7362.90 |
| 7029 | 0.0 | 1.0 | 29.60  | 346.45  |
| 7030 | 1.0 | 1.0 | 74.40  | 306.60  |
| 7031 | 1.0 | 1.0 | 105.65 | 6844.50 |

[7032 rows x 40 columns]

## 0.10 Feature Selection Using Correlation


```
[27]: #Get Correlation of "Churn" with other variables:
plt.figure(figsize=(15,8))
y_df = pd.DataFrame(y, columns=['Churn'])

# Concatenate the feature DataFrame and the target DataFrame
new_df = pd.concat([X_transformed_df, y_df], axis=1)

# Calculate correlations with the 'Churn' column
correlation = new_df.corr()['Churn'].sort_values(ascending=False)

# Plot the correlation of Churn with other variables
plt.figure(figsize=(15, 8))
correlation.plot(kind='bar')
plt.title("Correlation of Features with Churn")
plt.ylabel("Correlation Coefficient")
plt.show()
```

<Figure size 1500x800 with 0 Axes>



## 0.11 Feature Selection Using RandomForest

```
[28]: # Dataset has mixed values between numerical and categorical so its best to use
      ↪ randomforest rather than correlation matrix or chi-square

model = RandomForestClassifier()
model.fit(X_transformed, y)

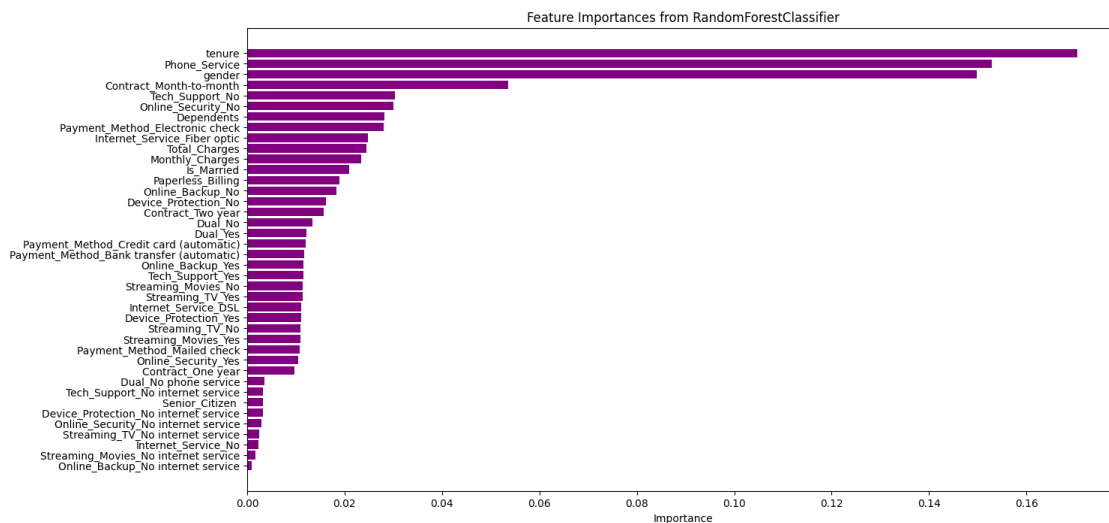
# Get feature importances
importances = model.feature_importances_
print(importances)

# Create a DataFrame to store feature names and their importances
features_df = pd.DataFrame({
    'Feature': X_transformed_df.columns,
    'Importance': importances
})

# Sort the features by importance in descending order
features_df = features_df.sort_values(by='Importance', ascending=True)
```

```
# Plot the feature importances
plt.figure(figsize=(15, 8))
plt.barh(features_df['Feature'], features_df['Importance'], color='purple')
plt.xlabel('Importance')
plt.title('Feature Importances from RandomForestClassifier')
plt.show()
```

```
[0.0134645  0.00362119 0.01211903 0.01114978 0.02472343 0.00236605
 0.02996051 0.00290799 0.01045103 0.0183466  0.00095171 0.0115077
 0.01623842 0.00317968 0.01109489 0.03035744 0.00331088 0.01148571
 0.01093526 0.00247654 0.01134598 0.01143882 0.00173175 0.01090508
 0.05363645 0.00967561 0.01567628 0.01177007 0.01208563 0.02799879
 0.01081188 0.02810681 0.02097648 0.02340859 0.01894302 0.15290328
 0.00322414 0.02444555 0.14982801 0.17043944]
```



## 0.12 Keeping High-Correlated Features Only

```
[29]: number_of_features = 15 # Take highest 15 features
filtered_features = []
for feature in features_df.tail(number_of_features).Feature:
    print(feature)
    filtered_features.append(feature)

for feature in X_transformed_df.columns:
    if feature not in filtered_features:
        X_transformed_df = X_transformed_df.drop(feature, axis=1)
```



```
print(X_transformed_df)
```

```
Online_Security_Yes
Payment_Method_Mailed check
Streaming_Movies_Yes
Streaming_TV_No
Device_Protection_Yes
Internet_Service_DSL
Streaming_TV_Yes
Streaming_Movies_No
Tech_Support_Yes
Online_Backup_Yes
Payment_Method_Bank transfer (automatic)
Payment_Method_Credit card (automatic)
Dual_Yes
Dual_No
Contract_Two year
Device_Protection_No
Online_Backup_No
Paperless_Billing
Is_Married
Monthly_Charges
Total_Charges
Internet_Service_Fiber optic
Payment_Method_Electronic check
Dependents
Online_Security_No
Tech_Support_No
Contract_Month-to-month
gender
Phone_Service
tenure
```

|      | Dual_No | Dual_Yes | Internet_Service_DSL | Internet_Service_Fiber optic | \ |
|------|---------|----------|----------------------|------------------------------|---|
| 0    | 0.0     | 0.0      | 1.0                  | 0.0                          |   |
| 1    | 1.0     | 0.0      | 1.0                  | 0.0                          |   |
| 2    | 1.0     | 0.0      | 1.0                  | 0.0                          |   |
| 3    | 0.0     | 0.0      | 1.0                  | 0.0                          |   |
| 4    | 1.0     | 0.0      | 0.0                  | 1.0                          |   |
| ...  | ...     | ...      | ...                  | ...                          |   |
| 7027 | 0.0     | 1.0      | 1.0                  | 0.0                          |   |
| 7028 | 0.0     | 1.0      | 0.0                  | 1.0                          |   |
| 7029 | 0.0     | 0.0      | 1.0                  | 0.0                          |   |
| 7030 | 0.0     | 1.0      | 0.0                  | 1.0                          |   |
| 7031 | 1.0     | 0.0      | 0.0                  | 1.0                          |   |

|   | Online_Security_No | Online_Security_Yes | Online_Backup_No | \ |
|---|--------------------|---------------------|------------------|---|
| 0 | 1.0                | 0.0                 | 0.0              |   |

|      |     |     |     |
|------|-----|-----|-----|
| 1    | 0.0 | 1.0 | 1.0 |
| 2    | 0.0 | 1.0 | 0.0 |
| 3    | 0.0 | 1.0 | 1.0 |
| 4    | 1.0 | 0.0 | 1.0 |
| ...  | ... | ... | ... |
| 7027 | 0.0 | 1.0 | 1.0 |
| 7028 | 1.0 | 0.0 | 0.0 |
| 7029 | 0.0 | 1.0 | 1.0 |
| 7030 | 1.0 | 0.0 | 1.0 |
| 7031 | 0.0 | 1.0 | 1.0 |

|      | Online_Backup_Yes | Device_Protection_No | Device_Protection_Yes | ... | \ |
|------|-------------------|----------------------|-----------------------|-----|---|
| 0    | 1.0               | 1.0                  | 0.0                   | ... |   |
| 1    | 0.0               | 0.0                  | 1.0                   | ... |   |
| 2    | 1.0               | 1.0                  | 0.0                   | ... |   |
| 3    | 0.0               | 0.0                  | 1.0                   | ... |   |
| 4    | 0.0               | 1.0                  | 0.0                   | ... |   |
| ...  | ...               | ...                  | ...                   | ... |   |
| 7027 | 0.0               | 0.0                  | 1.0                   | ... |   |
| 7028 | 1.0               | 0.0                  | 1.0                   | ... |   |
| 7029 | 0.0               | 1.0                  | 0.0                   | ... |   |
| 7030 | 0.0               | 1.0                  | 0.0                   | ... |   |
| 7031 | 0.0               | 0.0                  | 1.0                   | ... |   |

|      | Payment_Method_Electronic | check | Payment_Method_Mailed | check | \ |
|------|---------------------------|-------|-----------------------|-------|---|
| 0    |                           | 1.0   |                       | 0.0   |   |
| 1    |                           | 0.0   |                       | 1.0   |   |
| 2    |                           | 0.0   |                       | 1.0   |   |
| 3    |                           | 0.0   |                       | 0.0   |   |
| 4    |                           | 1.0   |                       | 0.0   |   |
| ...  |                           | ...   |                       | ...   |   |
| 7027 |                           | 0.0   |                       | 1.0   |   |
| 7028 |                           | 0.0   |                       | 0.0   |   |
| 7029 |                           | 1.0   |                       | 0.0   |   |
| 7030 |                           | 0.0   |                       | 1.0   |   |
| 7031 |                           | 0.0   |                       | 0.0   |   |

|      | Dependents | Is_Married | Monthly_Charges | Paperless_Billing | \ |
|------|------------|------------|-----------------|-------------------|---|
| 0    | 0.0        | 0.0        | 1.0             | 0.0               |   |
| 1    | 1.0        | 0.0        | 0.0             | 0.0               |   |
| 2    | 1.0        | 0.0        | 0.0             | 0.0               |   |
| 3    | 1.0        | 0.0        | 0.0             | 0.0               |   |
| 4    | 0.0        | 0.0        | 0.0             | 0.0               |   |
| ...  | ...        | ...        | ...             | ...               |   |
| 7027 | 1.0        | 0.0        | 1.0             | 1.0               |   |
| 7028 | 0.0        | 0.0        | 1.0             | 1.0               |   |
| 7029 | 0.0        | 0.0        | 1.0             | 1.0               |   |
| 7030 | 1.0        | 1.0        | 1.0             | 0.0               |   |

|      |     |     |     |     |
|------|-----|-----|-----|-----|
| 7031 | 1.0 | 0.0 | 0.0 | 0.0 |
|------|-----|-----|-----|-----|

|      | Phone_Service | Total_Charges | gender | tenure  |
|------|---------------|---------------|--------|---------|
| 0    | 1.0           | 1.0           | 29.85  | 29.85   |
| 1    | 34.0          | 0.0           | 56.95  | 1889.50 |
| 2    | 2.0           | 1.0           | 53.85  | 108.15  |
| 3    | 45.0          | 0.0           | 42.30  | 1840.75 |
| 4    | 2.0           | 1.0           | 70.70  | 151.65  |
| ...  | ...           | ...           | ...    | ...     |
| 7027 | 24.0          | 1.0           | 84.80  | 1990.50 |
| 7028 | 72.0          | 1.0           | 103.20 | 7362.90 |
| 7029 | 11.0          | 1.0           | 29.60  | 346.45  |
| 7030 | 4.0           | 1.0           | 74.40  | 306.60  |
| 7031 | 66.0          | 1.0           | 105.65 | 6844.50 |

[7032 rows x 30 columns]

## Splitting the dataset into the Training and Test sets

```
[30]: X_train, X_test, y_train, y_test = train_test_split(np.array(X_transformed_df),
↳ y, test_size=0.2, random_state=42)
```

### 0.13 Feature Scaling

```
[31]: sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

### 0.14 Perform UnderSampling

```
[32]: # from imblearn.under_sampling import RandomUnderSampler
# rus = RandomUnderSampler(random_state=42)
# X_train, y_train = rus.fit_resample(X_train, y_train)
```


### 0.15 Perform OverSampling

```
[33]: # from imblearn.over_sampling import SMOTE
# smote = SMOTE(random_state=42)
# X_train, y_train = smote.fit_resample(X_train, y_train)
```

### 0.16 Perform UnderOverSampling

```
[34]: from imblearn.combine import SMOTEENN
smote = SMOTEENN(sampling_strategy='auto', random_state=42)
X_train, y_train = smote.fit_resample(X_train, y_train)
```

```
[35]: pd.DataFrame(y_train).value_counts().plot.pie(autopct='%1.2f%%',
    ↪ colors=['skyblue', 'orange'])
plt.title('Class Distribution "Churn"')
plt.ylabel('')
plt.show()
```



## 0.17 Training XGBoost on the Training set

```
[36]: xgb_classifier = XGBClassifier(n_estimators=194, learning_rate=0.
    ↪ 0.08011727992081952, max_depth=3)
#{'n_estimators': 138, 'learning_rate': 0.2910194858724226, 'max_depth': 8}
xgb_classifier.fit(X_train, y_train)
```

```
[36]: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=0.08011727992081952,
    max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=3, max_leaves=None,
```

```
min_child_weight=None, missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=194, n_jobs=None,
num_parallel_tree=None, random_state=None, ...)
```

## 0.18 Confusion Matrix

```
[37]: y_pred = xgb_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[719 314]
 [ 80 294]]
```

```
[37]: 0.7199715707178393
```

## 0.19 K-Cross Validation

```
[38]: accuracies = cross_val_score(estimator = xgb_classifier, X = X_train, y =
    ↪ y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the
    ↪ variance
```

```
Accuracy: 95.65 %
```

```
Standard Deviation: 2.34 %
```

```
[39]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.70   | 0.78     | 1033    |
| 1            | 0.48      | 0.79   | 0.60     | 374     |
| accuracy     |           |        | 0.72     | 1407    |
| macro avg    | 0.69      | 0.74   | 0.69     | 1407    |
| weighted avg | 0.79      | 0.72   | 0.74     | 1407    |

```
[ ]: !pip install optuna
```

## 0.20 Hyper Parameter Tuning

```
[ ]: import optuna
# from sklearn.metrics import accuracy_score

def objective(trial):
    params = {
        'n_estimators': trial.suggest_int('n_estimators', 50, 200),
        'learning_rate': trial.suggest_loguniform('learning_rate', 0.01, 0.3),
        'max_depth': trial.suggest_int('max_depth', 3, 9),
    }
    model = XGBClassifier(**params)
    # model.fit(X_train, y_train)
    # Evaluate on the test set
    # test_accuracy = accuracy_score(y_test, model.predict(X_test))
    # return test_accuracy
    return cross_val_score(model, X_train, y_train, cv=5, scoring='accuracy').
    ↪mean()
study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=50)
print(study.best_params)
```

### 0.20.1 Trying Randomized Search

```
[42]: # param_dist = {
#     'learning_rate': [0.01, 0.1, 0.2, 0.3],
#     'n_estimators': [50, 100, 200],
#     'max_depth': [3, 5, 7, 9],
#     'min_child_weight': [1, 3, 5],
#     'subsample': [0.6, 0.8, 1.0]
# }

# random_search = RandomizedSearchCV(xgb_classifier,
    ↪param_distributions=param_dist, n_iter=10, cv=3)
# random_search.fit(X_train, y_train)
# best_accuracy = random_search.best_score_
# best_parameters = random_search.best_params_
# print(f"Best Accuracy: {round(best_accuracy*100, 2)} %")
# print(f"Best Parameters: {best_parameters}")
```

## 0.21 Training on CatBoost

```
[43]: cat_classifier = CatBoostClassifier(silent=True)
cat_classifier.fit(X_train, y_train)
```

```
[43]: <catboost.core.CatBoostClassifier at 0x7db287555e10>
```

```
[44]: y_pred = cat_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[732 301]
 [ 83 291]]
```

```
[44]: 0.7270788912579957
```

```
[45]: accuracies = cross_val_score(estimator = cat_classifier, X = X_train, y = y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the variance
```

Accuracy: 96.63 %

Standard Deviation: 1.93 %

```
[46]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.71   | 0.79     | 1033    |
| 1            | 0.49      | 0.78   | 0.60     | 374     |
| accuracy     |           |        | 0.73     | 1407    |
| macro avg    | 0.69      | 0.74   | 0.70     | 1407    |
| weighted avg | 0.79      | 0.73   | 0.74     | 1407    |

```
[47]: # import optuna
# def objective(trial):
#     # Suggest values for the hyperparameters
#     iterations = trial.suggest_categorical('iterations', [500, 1000])
#     learning_rate = trial.suggest_loguniform('learning_rate', 0.01, 0.2)
#     depth = trial.suggest_int('depth', 6, 10)
#     l2_leaf_reg = trial.suggest_float('l2_leaf_reg', 1, 5)

#     cat_classifier = CatBoostClassifier(
#         iterations=iterations,
#         learning_rate=learning_rate,
#         depth=depth,
#         l2_leaf_reg=l2_leaf_reg,
#         verbose=0
#     )
```

```
#     return cross_val_score(cat_classifier, X_train, y_train, cv=5,
    ↪scoring='accuracy').mean()
# study = optuna.create_study(direction='maximize')
# study.optimize(objective, n_trials=50)
# print(study.best_params)
```

```
[48]: # param_dist = {
#       'iterations': [500, 1000],
#       'learning_rate': [0.01, 0.1, 0.2],
#       'depth': [6, 8, 10],
#       'l2_leaf_reg': [1, 3, 5]
#     }

# random_search = RandomizedSearchCV(cat_classifier,
    ↪param_distributions=param_dist, n_iter=10, cv=3)
# random_search.fit(X_train, y_train)
# best_accuracy = random_search.best_score_
# best_parameters = random_search.best_params_
# print(f"Best Accuracy: {round(best_accuracy*100, 2)} %")
# print(f"Best Parameters: {best_parameters}")
```

## 0.22 Training KernelSVM on the Training set

```
[49]: svm_classifier = SVC(kernel = 'rbf', random_state = 0)
svm_classifier.fit(X_train, y_train)
```

```
[49]: SVC(random_state=0)
```

```
[50]: y_pred = svm_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[693 340]
 [ 80 294]]
```

```
[50]: 0.7014925373134329
```

```
[51]: from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = svm_classifier, X = X_train, y =
    ↪y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the
    ↪variance
```

```
Accuracy: 94.29 %
Standard Deviation: 0.9 %
```



```
[52]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.67   | 0.77     | 1033    |
| 1            | 0.46      | 0.79   | 0.58     | 374     |
| accuracy     |           |        | 0.70     | 1407    |
| macro avg    | 0.68      | 0.73   | 0.68     | 1407    |
| weighted avg | 0.78      | 0.70   | 0.72     | 1407    |

```
[53]: # from sklearn.model_selection import GridSearchCV
# parameters = [{'C': [0.25, 0.5, 0.75, 1], 'kernel' : ['linear']},
#               {'C': [0.25, 0.5, 0.75, 1], 'kernel' : ['rbf'], 'gamma' : [0.1,
#               ↳0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]}] # gamma parameter can only be used
#               ↳with the rbf kernel and not the linear one.
# grid_search = GridSearchCV(estimator = svm_classifier,
#                             param_grid = parameters,
#                             scoring = 'accuracy', # the metric with which you
#               ↳want to evaluate the performance of the model
#                             cv = 10, # number of trained test polls when
#               ↳applying k-fold cross validation
#                             n_jobs = -1) # number of cores to use
# grid_search.fit(X_train, y_train)
# best_accuracy = grid_search.best_score_
# best_parameters = grid_search.best_params_
# print(f"Best Accuracy: {round(best_accuracy*100, 2)} %")
# print(f"Best Parameters: {best_parameters}")
```

```
[54]: # def objective(trial):
#       # Suggest values for the hyperparameters
#       c = trial.suggest_loguniform('C', 1e-3, 1e3) # Loguniform for C to cover
#               ↳a wide range
#       kernel = trial.suggest_categorical('kernel', ['linear', 'poly', 'rbf',
#               ↳'sigmoid'])
#       degree = trial.suggest_int('degree', 2, 5) if kernel == 'poly' else 3 #
#               ↳Degree only relevant for 'poly'
#       gamma = trial.suggest_categorical('gamma', ['scale', 'auto'])
#
#       # Create an SVC with suggested parameters
#       svc = SVC(C=c, kernel=kernel, degree=degree, gamma=gamma)
#
#       # Perform cross-validation and return the mean accuracy
#       score = cross_val_score(svc, X_train, y_train, cv=5, scoring='accuracy').
#               ↳mean()
```

```
#         return score

# # Create and optimize the study
# study = optuna.create_study(direction='maximize')
# study.optimize(objective, n_trials=50)
# print(study.best_params)
```

## 0.23 Training Logistic Regression on the Training set

```
[55]: from sklearn.linear_model import LogisticRegression
lr_classifier = LogisticRegression(random_state = 0)
lr_classifier.fit(X_train, y_train)
```

```
[55]: LogisticRegression(random_state=0)
```

```
[56]: y_pred = lr_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[651 382]
 [ 54 320]]
```

```
[56]: 0.6901208244491827
```

```
[57]: from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = lr_classifier, X = X_train, y =
    ↪ y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the
    ↪ variance
```

```
Accuracy: 90.09 %
```

```
Standard Deviation: 1.16 %
```

```
[58]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.92      | 0.63   | 0.75     | 1033    |
| 1            | 0.46      | 0.86   | 0.59     | 374     |
| accuracy     |           |        | 0.69     | 1407    |
| macro avg    | 0.69      | 0.74   | 0.67     | 1407    |
| weighted avg | 0.80      | 0.69   | 0.71     | 1407    |

```
[59]: # def objective(trial):
#       # Suggest values for the hyperparameters
#       c = trial.suggest_loguniform('C', 1e-3, 1e3) # Regularization strength
#       ↪(log-uniform)
#       solver = trial.suggest_categorical('solver', ['liblinear', 'saga',
#       ↪'lbfgs', 'newton-cg'])
#       max_iter = trial.suggest_int('max_iter', 100, 500) # Maximum iterations
#       penalty = trial.suggest_categorical('penalty', ['l2', 'elasticnet']) #
#       ↪Penalty type
#
#       # ElasticNet mixing parameter (only relevant if 'penalty' is 'elasticnet')
#       l1_ratio = trial.suggest_float('l1_ratio', 0.0, 1.0) if penalty ==
#       ↪'elasticnet' else 0.5
#
#       # Create Logistic Regression with suggested parameters
#       logreg = LogisticRegression(
#           C=c,
#           solver=solver,
#           max_iter=max_iter,
#           penalty=penalty,
#           l1_ratio=l1_ratio,
#           multi_class='auto',
#           random_state=42
#       )
#
#       # Perform cross-validation and return the mean accuracy
#       score = cross_val_score(logreg, X_train, y_train, cv=5,
#       ↪scoring='accuracy').mean()
#       return score
#
# # Create and optimize the study
# study = optuna.create_study(direction='maximize')
# study.optimize(objective, n_trials=50)
# print(study.best_params)
```

## 0.24 Training Random Forest on the Training set

```
[60]: from sklearn.ensemble import RandomForestClassifier
rf_classifier = RandomForestClassifier(n_estimators= 220, max_depth=17,
↪criterion= 'entropy', random_state= 0, min_samples_split=5,
↪min_samples_leaf=1, max_features='sqrt')
rf_classifier.fit(X_train, y_train)
# {'n_estimators': 220, 'max_depth': 17, 'min_samples_split': 5,
↪'min_samples_leaf': 1, 'max_features': 'sqrt'}
```

```
[60]: RandomForestClassifier(criterion='entropy', max_depth=17, min_samples_split=5,
                             n_estimators=220, random_state=0)
```

```
[61]: y_pred = rf_classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
[[717 316]
 [ 81 293]]
```

```
[61]: 0.7178393745557925
```

```
[62]: from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = rf_classifier, X = X_train, y =
    y_train, cv = 10) # number of folds
print(f"Accuracy: {round(accuracies.mean()*100, 2)} %")
print(f"Standard Deviation: {round(accuracies.std()*100, 2)} %") # for the
    variance
```

Accuracy: 96.43 %

Standard Deviation: 1.21 %

```
[63]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.69   | 0.78     | 1033    |
| 1            | 0.48      | 0.78   | 0.60     | 374     |
| accuracy     |           |        | 0.72     | 1407    |
| macro avg    | 0.69      | 0.74   | 0.69     | 1407    |
| weighted avg | 0.79      | 0.72   | 0.73     | 1407    |

```
[64]: # import optuna
# #{'n_estimators': 220, 'max_depth': 17, 'min_samples_split': 5,
    # 'min_samples_leaf': 1, 'max_features': 'sqrt'}
# def objective(trial):
#     # Suggest values for the hyperparameters
#     n_estimators = trial.suggest_int('n_estimators', 50, 500) # Number of
    # trees
#     max_depth = trial.suggest_int('max_depth', 2, 20) # Maximum depth of
    # trees
#     min_samples_split = trial.suggest_int('min_samples_split', 2, 20) #
    # Minimum samples to split a node
```

```

#     min_samples_leaf = trial.suggest_int('min_samples_leaf', 1, 20) #
    ↪ Minimum samples per leaf
#     max_features = trial.suggest_categorical('max_features', ['sqrt', 'log2',
    ↪ None]) # Number of features to consider for splits

#     # Create a RandomForestClassifier with suggested parameters
#     rf = RandomForestClassifier(
#         n_estimators=n_estimators,
#         max_depth=max_depth,
#         min_samples_split=min_samples_split,
#         min_samples_leaf=min_samples_leaf,
#         max_features=max_features,
#         random_state=42
#     )

#     # Perform cross-validation and return the mean accuracy
#     score = cross_val_score(rf, X_train, y_train, cv=5, scoring='accuracy').
    ↪ mean()
#     return score

# # Create and optimize the study
# study = optuna.create_study(direction='maximize')
# study.optimize(objective, n_trials=50)
# print(study.best_params)

```

### 0.24.1 Building an ANN

```

[76]: import tensorflow as tf
from tensorflow.keras.regularizers import l2
# Building the Network

model = tf.keras.models.Sequential(
    [
        tf.keras.Input(shape=(X_train.shape[1],)),
        tf.keras.layers.Dense(units=256, activation='relu'),
        tf.keras.layers.BatchNormalization(),
        tf.keras.layers.Dropout(0.3),
        tf.keras.layers.Dense(units=256, activation='relu',
    ↪ kernel_regularizer=l2(0.001)),
        tf.keras.layers.Dropout(0.2),
        tf.keras.layers.Dense(units=128, activation='relu'),
        tf.keras.layers.Dense(units=1, activation='sigmoid')
    ]
)

```

```
[77]: def plot_loss_acc(history):
    '''Plots the training and validation loss and accuracy from a history_
    ↪object'''
    acc = history.history['accuracy']
    loss = history.history['loss']
    val_acc = history.history['val_accuracy']
    val_loss = history.history['val_loss']

    epochs = range(len(acc))

    fig, ax = plt.subplots(1,2, figsize=(12, 6))
    ax[0].plot(epochs, acc, 'b', label='Training accuracy')
    ax[0].plot(epochs, val_acc, 'r', label='Validation accuracy')
    ax[0].set_title('Training accuracy')
    ax[0].set_xlabel('epochs')
    ax[0].set_ylabel('accuracy')
    ax[0].legend()

    ax[1].plot(epochs, loss, 'b', label='Training Loss')
    ax[1].plot(epochs, val_loss, 'r', label='Validation Loss')
    ax[1].set_title('Training loss')
    ax[1].set_xlabel('epochs')
    ax[1].set_ylabel('loss')
    ax[1].legend()

    plt.show()
```

```
[78]: class EarlyStoppingCallback(tf.keras.callbacks.Callback):

    # Define the correct function signature for on_epoch_end method
    def on_epoch_end(self, epoch, logs={}):

        # Check if the accuracy is greater or equal to 0.98
        if (logs.get("accuracy") >= 0.94):

            # Stop training once the above condition is met
            self.model.stop_training = True

            print("\nReached 94% accuracy so cancelling training!")
```

```
[79]: def plot_learning_rate(history):
    """
    Plot learning rate vs loss.
    """

    # Define the learning rate array
    lrs = 1e-8 * (10 ** (np.arange(len(history.history["loss"])) / 20))
```

```

# Set the figure size
plt.figure(figsize=(10, 6))

# Set the grid
plt.grid(True)

# Plot the loss in log scale
plt.semilogx(lrs, history.history["loss"])

# Increase the tickmarks size
plt.tick_params('both', length=10, width=1, which='both')

# Set the plot boundaries
plt.axis([1e-8, 8e-4, 0, 50])
plt.ylim(0, 1.0) # This would zoom the y-axis between 0 and 1.0

# Draw the graph on screen
plt.xlabel('Learning Rate')
plt.ylabel('Loss')
plt.title('Learning Rate vs Loss')
plt.show()

```

```

[80]: def tuning_learn_rate(model, X_train, y_train, optimizer, loss, batch_size=64,
    ↪ epochs=100):
    """
    Tunes the learning rate using an exponential increase.

    Args:
        model (tf.keras.Model): The Keras model.
        X_train (np.array): Training features.
        y_train (np.array): Training labels.
        optimizer (tf.keras.optimizers): Optimizer to use.
        loss (str or tf.keras.losses.Loss): Loss function for the model.
        epochs (int): Number of epochs to tune learning rate.
        batch_size (int): The batch size for training.

    Returns:
        history (tf.keras.callbacks.History): Training history object.
    """
    # Convert the data into a TensorFlow Dataset
    train_dataset = tf.data.Dataset.from_tensor_slices((X_train, y_train))
    train_dataset = train_dataset.batch(batch_size)

    # Define the learning rate scheduler (exponential increase)
    lr_scheduler = tf.keras.callbacks.LearningRateScheduler(
        lambda epoch: 1e-8 * 10**(epoch / 20) # Exponential growth

```

```

)

# Compile the model with the optimizer and loss function
model.compile(
    optimizer=optimizer,
    loss=loss
)

# Fit the model with the learning rate scheduler
history = model.fit(train_dataset, epochs=epochs, callbacks=[lr_scheduler])
return history

# Example usage
tuned_history = tuning_learn_rate(
    model, X_train, y_train, optimizer=tf.keras.optimizers.Adam(),
    ↪loss='binary_crossentropy'
)

```

```

Epoch 1/100
80/80          6s 13ms/step -
loss: 1.0233 - learning_rate: 1.0000e-08
Epoch 2/100
80/80          1s 13ms/step -
loss: 1.0207 - learning_rate: 1.1220e-08
Epoch 3/100
80/80          1s 11ms/step -
loss: 1.0322 - learning_rate: 1.2589e-08
Epoch 4/100
80/80          1s 10ms/step -
loss: 1.0249 - learning_rate: 1.4125e-08
Epoch 5/100
80/80          1s 11ms/step -
loss: 1.0232 - learning_rate: 1.5849e-08
Epoch 6/100
80/80          1s 11ms/step -
loss: 1.0188 - learning_rate: 1.7783e-08
Epoch 7/100
80/80          1s 10ms/step -
loss: 1.0324 - learning_rate: 1.9953e-08
Epoch 8/100
80/80          1s 11ms/step -
loss: 1.0155 - learning_rate: 2.2387e-08
Epoch 9/100
80/80          1s 10ms/step -
loss: 1.0331 - learning_rate: 2.5119e-08
Epoch 10/100
80/80          1s 10ms/step -

```



```

loss: 1.0210 - learning_rate: 2.8184e-08
Epoch 11/100
80/80          1s 10ms/step -
loss: 1.0219 - learning_rate: 3.1623e-08
Epoch 12/100
80/80          2s 14ms/step -
loss: 1.0172 - learning_rate: 3.5481e-08
Epoch 13/100
80/80          1s 14ms/step -
loss: 1.0205 - learning_rate: 3.9811e-08
Epoch 14/100
80/80          1s 15ms/step -
loss: 1.0282 - learning_rate: 4.4668e-08
Epoch 15/100
80/80          1s 15ms/step -
loss: 1.0297 - learning_rate: 5.0119e-08
Epoch 16/100
80/80          1s 8ms/step - loss:
1.0256 - learning_rate: 5.6234e-08
Epoch 17/100
80/80          1s 11ms/step -
loss: 1.0223 - learning_rate: 6.3096e-08
Epoch 18/100
80/80          1s 11ms/step -
loss: 1.0132 - learning_rate: 7.0795e-08
Epoch 19/100
80/80          1s 5ms/step - loss:
1.0321 - learning_rate: 7.9433e-08
Epoch 20/100
80/80          1s 6ms/step - loss:
1.0366 - learning_rate: 8.9125e-08
Epoch 21/100
80/80          1s 5ms/step - loss:
1.0214 - learning_rate: 1.0000e-07
Epoch 22/100
80/80          0s 6ms/step - loss:
1.0249 - learning_rate: 1.1220e-07
Epoch 23/100
80/80          0s 5ms/step - loss:
1.0139 - learning_rate: 1.2589e-07
Epoch 24/100
80/80          1s 5ms/step - loss:
1.0235 - learning_rate: 1.4125e-07
Epoch 25/100
80/80          1s 6ms/step - loss:
1.0243 - learning_rate: 1.5849e-07
Epoch 26/100
80/80          0s 5ms/step - loss:

```

1.0291 - learning\_rate: 1.7783e-07  
Epoch 27/100  
80/80 1s 6ms/step - loss:  
1.0254 - learning\_rate: 1.9953e-07  
Epoch 28/100  
80/80 1s 5ms/step - loss:  
1.0145 - learning\_rate: 2.2387e-07  
Epoch 29/100  
80/80 1s 5ms/step - loss:  
1.0148 - learning\_rate: 2.5119e-07  
Epoch 30/100  
80/80 1s 8ms/step - loss:  
1.0196 - learning\_rate: 2.8184e-07  
Epoch 31/100  
80/80 1s 10ms/step -  
loss: 1.0243 - learning\_rate: 3.1623e-07  
Epoch 32/100  
80/80 1s 10ms/step -  
loss: 1.0250 - learning\_rate: 3.5481e-07  
Epoch 33/100  
80/80 1s 5ms/step - loss:  
1.0286 - learning\_rate: 3.9811e-07  
Epoch 34/100  
80/80 1s 6ms/step - loss:  
1.0345 - learning\_rate: 4.4668e-07  
Epoch 35/100  
80/80 1s 6ms/step - loss:  
1.0233 - learning\_rate: 5.0119e-07  
Epoch 36/100  
80/80 1s 5ms/step - loss:  
1.0232 - learning\_rate: 5.6234e-07  
Epoch 37/100  
80/80 1s 5ms/step - loss:  
1.0181 - learning\_rate: 6.3096e-07  
Epoch 38/100  
80/80 0s 6ms/step - loss:  
1.0299 - learning\_rate: 7.0795e-07  
Epoch 39/100  
80/80 0s 6ms/step - loss:  
1.0230 - learning\_rate: 7.9433e-07  
Epoch 40/100  
80/80 0s 6ms/step - loss:  
1.0220 - learning\_rate: 8.9125e-07  
Epoch 41/100  
80/80 0s 6ms/step - loss:  
1.0297 - learning\_rate: 1.0000e-06  
Epoch 42/100  
80/80 1s 6ms/step - loss:

1.0267 - learning\_rate: 1.1220e-06  
Epoch 43/100  
80/80 1s 6ms/step - loss:  
1.0297 - learning\_rate: 1.2589e-06  
Epoch 44/100  
80/80 1s 5ms/step - loss:  
1.0242 - learning\_rate: 1.4125e-06  
Epoch 45/100  
80/80 0s 6ms/step - loss:  
1.0324 - learning\_rate: 1.5849e-06  
Epoch 46/100  
80/80 1s 5ms/step - loss:  
1.0323 - learning\_rate: 1.7783e-06  
Epoch 47/100  
80/80 0s 6ms/step - loss:  
1.0356 - learning\_rate: 1.9953e-06  
Epoch 48/100  
80/80 0s 5ms/step - loss:  
1.0256 - learning\_rate: 2.2387e-06  
Epoch 49/100  
80/80 1s 6ms/step - loss:  
1.0206 - learning\_rate: 2.5119e-06  
Epoch 50/100  
80/80 1s 5ms/step - loss:  
1.0333 - learning\_rate: 2.8184e-06  
Epoch 51/100  
80/80 1s 10ms/step -  
loss: 1.0158 - learning\_rate: 3.1623e-06  
Epoch 52/100  
80/80 1s 10ms/step -  
loss: 1.0163 - learning\_rate: 3.5481e-06  
Epoch 53/100  
80/80 1s 8ms/step - loss:  
1.0304 - learning\_rate: 3.9811e-06  
Epoch 54/100  
80/80 0s 5ms/step - loss:  
1.0158 - learning\_rate: 4.4668e-06  
Epoch 55/100  
80/80 1s 6ms/step - loss:  
1.0039 - learning\_rate: 5.0119e-06  
Epoch 56/100  
80/80 1s 5ms/step - loss:  
1.0127 - learning\_rate: 5.6234e-06  
Epoch 57/100  
80/80 1s 5ms/step - loss:  
0.9992 - learning\_rate: 6.3096e-06  
Epoch 58/100  
80/80 1s 6ms/step - loss:

```

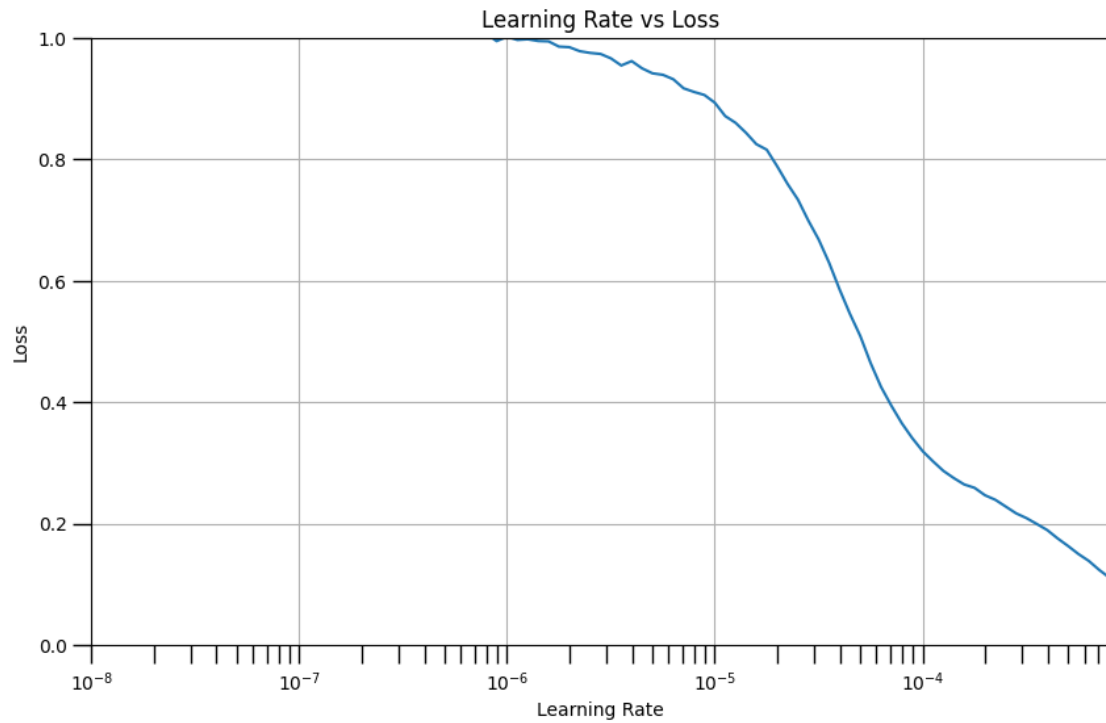
0.9882 - learning_rate: 7.0795e-06
Epoch 59/100
80/80          0s 5ms/step - loss:
0.9902 - learning_rate: 7.9433e-06
Epoch 60/100
80/80          1s 6ms/step - loss:
0.9839 - learning_rate: 8.9125e-06
Epoch 61/100
80/80          0s 5ms/step - loss:
0.9709 - learning_rate: 1.0000e-05
Epoch 62/100
80/80          1s 6ms/step - loss:
0.9464 - learning_rate: 1.1220e-05
Epoch 63/100
80/80          1s 5ms/step - loss:
0.9408 - learning_rate: 1.2589e-05
Epoch 64/100
80/80          1s 5ms/step - loss:
0.9237 - learning_rate: 1.4125e-05
Epoch 65/100
80/80          1s 6ms/step - loss:
0.9007 - learning_rate: 1.5849e-05
Epoch 66/100
80/80          1s 6ms/step - loss:
0.9043 - learning_rate: 1.7783e-05
Epoch 67/100
80/80          1s 6ms/step - loss:
0.8718 - learning_rate: 1.9953e-05
Epoch 68/100
80/80          1s 6ms/step - loss:
0.8349 - learning_rate: 2.2387e-05
Epoch 69/100
80/80          1s 6ms/step - loss:
0.8053 - learning_rate: 2.5119e-05
Epoch 70/100
80/80          1s 8ms/step - loss:
0.7787 - learning_rate: 2.8184e-05
Epoch 71/100
80/80          1s 10ms/step -
loss: 0.7432 - learning_rate: 3.1623e-05
Epoch 72/100
80/80          1s 10ms/step -
loss: 0.7056 - learning_rate: 3.5481e-05
Epoch 73/100
80/80          1s 11ms/step -
loss: 0.6511 - learning_rate: 3.9811e-05
Epoch 74/100
80/80          1s 6ms/step - loss:

```

0.6035 - learning\_rate: 4.4668e-05  
Epoch 75/100  
80/80 1s 5ms/step - loss:  
0.5657 - learning\_rate: 5.0119e-05  
Epoch 76/100  
80/80 0s 6ms/step - loss:  
0.5150 - learning\_rate: 5.6234e-05  
Epoch 77/100  
80/80 1s 6ms/step - loss:  
0.4717 - learning\_rate: 6.3096e-05  
Epoch 78/100  
80/80 0s 5ms/step - loss:  
0.4254 - learning\_rate: 7.0795e-05  
Epoch 79/100  
80/80 0s 5ms/step - loss:  
0.4001 - learning\_rate: 7.9433e-05  
Epoch 80/100  
80/80 1s 6ms/step - loss:  
0.3654 - learning\_rate: 8.9125e-05  
Epoch 81/100  
80/80 0s 6ms/step - loss:  
0.3415 - learning\_rate: 1.0000e-04  
Epoch 82/100  
80/80 0s 6ms/step - loss:  
0.3244 - learning\_rate: 1.1220e-04  
Epoch 83/100  
80/80 1s 6ms/step - loss:  
0.3043 - learning\_rate: 1.2589e-04  
Epoch 84/100  
80/80 1s 5ms/step - loss:  
0.2898 - learning\_rate: 1.4125e-04  
Epoch 85/100  
80/80 1s 6ms/step - loss:  
0.2785 - learning\_rate: 1.5849e-04  
Epoch 86/100  
80/80 1s 6ms/step - loss:  
0.2744 - learning\_rate: 1.7783e-04  
Epoch 87/100  
80/80 0s 6ms/step - loss:  
0.2557 - learning\_rate: 1.9953e-04  
Epoch 88/100  
80/80 1s 6ms/step - loss:  
0.2479 - learning\_rate: 2.2387e-04  
Epoch 89/100  
80/80 0s 5ms/step - loss:  
0.2384 - learning\_rate: 2.5119e-04  
Epoch 90/100  
80/80 0s 6ms/step - loss:

```
0.2238 - learning_rate: 2.8184e-04
Epoch 91/100
80/80          1s 5ms/step - loss:
0.2172 - learning_rate: 3.1623e-04
Epoch 92/100
80/80          1s 10ms/step -
loss: 0.2081 - learning_rate: 3.5481e-04
Epoch 93/100
80/80          1s 10ms/step -
loss: 0.1982 - learning_rate: 3.9811e-04
Epoch 94/100
80/80          1s 11ms/step -
loss: 0.1848 - learning_rate: 4.4668e-04
Epoch 95/100
80/80          1s 11ms/step -
loss: 0.1725 - learning_rate: 5.0119e-04
Epoch 96/100
80/80          0s 6ms/step - loss:
0.1585 - learning_rate: 5.6234e-04
Epoch 97/100
80/80          0s 6ms/step - loss:
0.1497 - learning_rate: 6.3096e-04
Epoch 98/100
80/80          1s 6ms/step - loss:
0.1313 - learning_rate: 7.0795e-04
Epoch 99/100
80/80          0s 5ms/step - loss:
0.1191 - learning_rate: 7.9433e-04
Epoch 100/100
80/80          0s 6ms/step - loss:
0.1082 - learning_rate: 8.9125e-04
```

```
[81]: plot_learning_rate(tuned_history)
```



[82]: *# Compiling the Network*

```
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=2e-4),  
              ↪ loss='binary_crossentropy', metrics=['accuracy'])  
# learning_rate=4e-4
```

[83]: model.summary()

Model: "sequential\_2"

| Layer (type)          | Output Shape |   |
|-----------------------|--------------|---|
| ↪ Param #             |              |   |
| dense_8 (Dense)       | (None, 256)  | ↪ |
| ↪ 7,936               |              |   |
| batch_normalization_2 | (None, 256)  | ↪ |
| ↪ 1,024               |              |   |
| (BatchNormalization)  |              | ↪ |
| ↪                     |              |   |

```

dropout_4 (Dropout)                (None, 256)
↳ 0
dense_9 (Dense)                    (None, 256)
↳ 65,792
dropout_5 (Dropout)                (None, 256)
↳ 0
dense_10 (Dense)                   (None, 128)
↳ 32,896
dense_11 (Dense)                   (None, 1)
↳ 129

```

Total params: 107,777 (421.00 KB)

Trainable params: 107,265 (419.00 KB)

Non-trainable params: 512 (2.00 KB)

```

[84]: # Training the data on the network

history = model.fit(x=X_train, y=y_train, validation_split=0.2, batch_size=32,
↳ epochs=50, callbacks=[EarlyStoppingCallback()])

```

```

Epoch 1/50
127/127      3s 7ms/step -
accuracy: 0.7858 - loss: 0.9670 - val_accuracy: 0.8571 - val_loss: 0.5272
Epoch 2/50
127/127      1s 5ms/step -
accuracy: 0.8886 - loss: 0.4006 - val_accuracy: 0.8512 - val_loss: 0.4961
Epoch 3/50
127/127      1s 5ms/step -
accuracy: 0.8983 - loss: 0.3592 - val_accuracy: 0.8897 - val_loss: 0.3633
Epoch 4/50
127/127      2s 8ms/step -
accuracy: 0.9018 - loss: 0.3097 - val_accuracy: 0.8916 - val_loss: 0.3459
Epoch 5/50
127/127      1s 8ms/step -
accuracy: 0.9086 - loss: 0.2991 - val_accuracy: 0.9005 - val_loss: 0.3372
Epoch 6/50
127/127      1s 9ms/step -
accuracy: 0.9114 - loss: 0.2835 - val_accuracy: 0.9005 - val_loss: 0.3215

```



```

Epoch 7/50
127/127          1s 9ms/step -
accuracy: 0.9239 - loss: 0.2501 - val_accuracy: 0.9015 - val_loss: 0.3110
Epoch 8/50
127/127          1s 6ms/step -
accuracy: 0.9322 - loss: 0.2507 - val_accuracy: 0.9084 - val_loss: 0.2988
Epoch 9/50
127/127          1s 5ms/step -
accuracy: 0.9279 - loss: 0.2456 - val_accuracy: 0.9064 - val_loss: 0.2903
Epoch 10/50
127/127          1s 5ms/step -
accuracy: 0.9327 - loss: 0.2319 - val_accuracy: 0.8926 - val_loss: 0.3139
Epoch 11/50
127/127          1s 5ms/step -
accuracy: 0.9371 - loss: 0.2196 - val_accuracy: 0.9241 - val_loss: 0.2346
Epoch 12/50
124/127          0s 4ms/step -
accuracy: 0.9379 - loss: 0.2258
Reached 94% accuracy so cancelling training!
127/127          1s 5ms/step -
accuracy: 0.9380 - loss: 0.2256 - val_accuracy: 0.9103 - val_loss: 0.2813

```

```
[85]: plot_loss_acc(history)
```



```

[86]: y_pred = model.predict(X_test)
      y_pred = (y_pred > 0.5)
      print(np.concatenate((y_pred.reshape(-1, 1), y_test.reshape(-1, 1)), 1))

```

44/44 0s 3ms/step

```
[[0 0]
 [0 0]
 [1 1]
 ...
 [0 0]
 [0 0]
 [0 0]]
```

```
[87]: from sklearn.metrics import confusion_matrix, accuracy_score
      cm = confusion_matrix(y_test, y_pred)
      print(cm)
      print(accuracy_score(y_test, y_pred))
```

```
[[735 298]
 [ 96 278]]
0.7199715707178393
```

```
[88]: from sklearn.metrics import classification_report
      print(classification_report(y_test, y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.88      | 0.71   | 0.79     | 1033    |
| 1            | 0.48      | 0.74   | 0.59     | 374     |
| accuracy     |           |        | 0.72     | 1407    |
| macro avg    | 0.68      | 0.73   | 0.69     | 1407    |
| weighted avg | 0.78      | 0.72   | 0.73     | 1407    |